

zkID technical report

Jane Doe^{1,2} and John Doe¹

¹ Institute A, City, Country, jane@institute

² Institute B, City, Country, john@institute

Abstract. Main deliveries: 1. Technical report on zk component for the digital id wallet 2. A comparison with current works 3. Applying to EUDI.

Keywords: Anonymous credential · programmable zkp

1 Introduction

People use digital credentials to prove a single fact—such as being over 18, holding a license, or belonging to an organization—without presenting an entire document. In the W3C model [Wor25], an issuer signs a credential, a holder keeps it (typically in a wallet application), and a verifier checks a presentation—that is, data derived from one or more credentials and shown to a specific verifier for a specific session. We take this as the baseline flow.

A practical privacy risk arises when the same credential is shown to different services over time. If a presentation exposes a stable technical trace—for example, an identifier or handle that repeats—then services can compare logs later and infer that the sessions likely came from the same source. We call this linkability: separate presentations that should remain independent become easy to connect. Standards documents explicitly caution against this: the W3C Verifiable Credentials Data Model states that securing mechanisms must not leak information that would enable a verifier to correlate a holder across multiple presentations [Wor25]; and NIST’s federation guidance recommends pairwise pseudonymous identifiers—meaning unique, opaque identifiers for each relying party—to ensure separation across services and reduce the risk of user tracking [GRS⁺]. Deployments typically provide two recurring capabilities: device binding (confirming that the credential is presented by the authorized wallet) and revocation/status (learning whether the credential is still valid). To support these, systems often expose stable values. For device binding, some ecosystems tie a credential to a long-lived cryptographic key under the control of the device so a verifier can accept only “the right wallet”; if the same key reference or its proof is visible across sessions, different verifiers can match their logs [FYC25]. For status, verifiers query whether a credential was suspended or revoked; a one-to-one status reference (for example, a unique status URL per credential) lets the status provider observe which verifier asked about which holder and when [Wor25]. These two mechanisms keep systems reliable and usable, but because they rely on fixed values, they also leave repeating patterns that make it possible to link separate credential uses.

Device binding and revocation introduce fixed points outside the credential itself: the first in verifier transcripts, the second in status queries. When such points remain stable across uses, separate sessions can be matched. In what follows, we compare designs using a simple reference view: issuers and credential formats remain as deployed; privacy logic is handled on the holder side; and presentations are per-session and non-repeating. We use this view to organize related work next.

1.1 Related Work

For the purpose of comparison, we first outline a reference architecture that represents what an anonymous-credential system would ideally look like if it is to integrate smoothly with current infrastructures. In this model, the issuer is treated as fixed components that continue to use their existing public-key algorithms (such as RSA or ECDSA) and standard credential formats (e.g., JWT or mDL), since it's typically difficult to change once deployed. All additional logic is placed in the user's wallet and the verifier. The wallet is expected to operate in two stages: an offline Prepare step, which verifies the issuer's signature once using standard libraries, parses and normalizes credential attributes (for example, turning a date of birth into an integer age), and commits to those attributes using a binding and hiding commitment scheme (a cryptographic way to lock values so they can later be revealed or proven in restricted form); and an online Show step, which runs per presentation, where the wallet selects only the attributes or predicates required by a relying party's policy, proves them in zero knowledge against the stored commitments, and includes a fresh device signature over the session challenge to ensure the proof is tied to the holder's device. A further requirement is modularity: each major function—issuer signature verification, attribute commitment, predicate proofs, and device binding—should be defined as a separate module with a clear interface. This separation makes it possible to swap the underlying proof engine (for example, using a SNARK today or a post-quantum proof system in the future) without requiring changes to parts of the system that are costly or impractical to modify. The purpose of this modular view is to act as a comparison framework: it outlines how a deployment-friendly anonymous-credential stack could be structured, making it easier to compare proposals by the modules they cover, the constraints they address, and the trade-offs they make.

Anonymous Credentials from ECDSA. [Fas24] This work considers environments where credential issuers already sign with ECDSA on standardized curves (such as P-256) and hash data with SHA-256. The main challenge is that proving correctness of an ECDSA signature in zero knowledge is costly with standard proof systems, because the arithmetic used in P-256 and the bit-level operations in SHA-256 do not align well with the fast polynomial techniques (such as number-theoretic transforms, a method that speeds up polynomial multiplication over special fields) that many modern ZK libraries rely on. To handle this, the authors introduce custom circuits for ECDSA and SHA-256, and use a layered protocol based on the sum-check technique and a lightweight encoding (Reed-Solomon code) to control proof size. An additional “consistency check” ensures that the same hidden signing key is used across both the signature and the hash logic. At presentation, the wallet produces a proof for the verifier and the device also signs a fresh challenge (this is the device-binding step: a live signature that ties the proof to the holder's device). In terms of the reference system view, issuer compatibility is preserved, selective disclosure is supported, and device binding is included; however, there is no reusable offline phase, so the full proof is generated at every presentation. The reported costs are about 60 ms to prove one ECDSA signature and about 1.2 s for a complete mDL presentation on mobile devices [Fas24, §5.3, §6.2], with larger proof sizes and higher verifier effort than systems based on succinct setup-dependent SNARKs.

Crescent Credentials. [PPZ24] This work considers environments where issuers continue using existing credential formats such as JWT or mDL and their current signing keys, so no issuer-side changes are required. Its workflow is split into a heavy one-time Prepare phase and a lightweight per-presentation Show phase. In Prepare, the wallet verifies the issuer's signature, parses the credential into attributes, and creates two reusable artifacts—that is, cryptographic objects the wallet reuses across presentations: (i) a Groth16 proof that these checks were done correctly, and (ii) a Pedersen vector commitment over the attributes,

enabling selective disclosure. Both artifacts support re-randomization for unlinkability. In the Show phase, the wallet re-randomizes the prepared artifacts and attaches only the proofs required by the verifier’s policy, such as proving an age threshold or linking two credentials to the same holder. Device binding can be added at this step by letting the secure element sign the verifier’s challenge. In terms of the reference system view, Crescent realizes the two-phase design with reusable offline work and modular predicates, while leaving issuers unchanged. The trade-offs are significant: the Prepare phase is heavy (tens of seconds for JWTs and minutes for mDLs), the scheme depends on pairing-based Groth16 proofs with a large universal setup (≈ 661 MB–1.1 GB [PPZ24, §4]), and the security model is classical only, without post-quantum protection. The Show step, however, runs with low latency—typically 22–41 ms with ≈ 1 KB proofs, or about 315 ms with device binding [PPZ24, §4].

zkID. Our construction works with standardized credentials (e.g., SD-JWT, mDL) and existing PKI (RSA/ECDSA), so issuers do not need to change their issuance pipelines. The zkID workflow follows the two-phase split in the reference view: a one-time Prepare phase and a per-presentation Show phase. In Prepare, the wallet verifies the issuer’s signature, parses the credential into normalized messages, computes the associated hashes, and produces two reusable artifacts: (i) zero-knowledge proofs that issuer-side checks and parsing were done correctly, and (ii) Hyrax-style Pedersen vector commitments to a designated message column, supporting efficient proofs over multiple attributes. In Show, the wallet proves only the verifier’s requested predicates and includes a fresh device-binding signature. To link Prepare and Show without revealing values, the verifier checks equality of commitments across both proofs; the wallet reuses the corresponding randomness for that session. The proving backend is transparent (no trusted setup). It checks the arithmetic constraints with a sum-check-style protocol and uses a small inner-product check to verify commitment openings. For device binding, we choose a curve whose scalar field matches the device’s signature field (e.g., P-256), so the device signature can be verified directly inside the proof without emulation or field translation. In terms of the reference system view, issuer compatibility is preserved, the two-phase reuse is integrated into the workflow, predicates are modular, and there is no trusted setup. The trade-offs are that security currently relies on discrete-log assumptions (not post-quantum) and that commitment equality requires using the same curve across Prepare and Show; the modular interface leaves room to swap in lattice-based commitments when suitable.

Taken together, these approaches show the current trade-off: systems either reuse existing issuer infrastructure but pay high per-presentation costs, or they achieve fast online proofs at the price of large setups and pairing-based assumptions. Our construction, zkID, aims to combine issuer compatibility with reusable offline work, while remaining transparent and modular. Before presenting the design, we review the algebraic and cryptographic preliminaries that will be used throughout the paper.

2 Preliminaries - WIP

Notation For $n \in \mathbb{N}$ we write $[n] = \{1, \dots, n\}$. Bold letters denote vectors, e.g., $\mathbf{m} = (m_1, \dots, m_n)$. Concatenation is written \parallel . The security parameter is λ ; $\text{negl}(\lambda)$ denotes a negligible function. For a (possibly randomized) algorithm Alg , we write $y \leftarrow \text{Alg}(x)$ for its output on input x .

There are three roles:

- The *issuer* I signs credentials with a long-term key pair (SK_I, PK_I) (e.g., ECDSA P-256 or RSA).
- The *prover* P is the holder’s wallet, which stores credentials and generates proofs.

- The *verifier* V is the relying party that checks proofs against a policy.

For device binding, the prover’s secure element holds an additional signing key pair (SK_D, PK_D) used only to sign fresh per-session challenges.

Credentials A credential is a standardized signed object S (e.g., SD-JWT [FYC25] or mDL [fS21]). Parsing maps S into an ordered vector of attributes

$$\mathbf{m} = (m_1, \dots, m_n).$$

Non-numeric fields (strings, dates) are encoded injectively into integers. The resulting integers are interpreted in a prime field $\mathbb{F} = \mathbb{F}_q$ chosen for the proof backend. For each attribute m_i we sample a salt $s_i \leftarrow \mathbb{F}$ and compute

$$h_i = H(m_i \parallel s_i),$$

where H is instantiated as SHA-256. The issuer’s signature is

$$\sigma_I = \text{Sign}_{SK_I}(h_1, \dots, h_n),$$

verified under PK_I .

Commitments and Proof Interface To support selective disclosure without revealing raw attributes, the wallet commits to \mathbf{m} using Pedersen vector commitments. Let \mathbb{G} be a cyclic group of prime order q with public generators (g_1, \dots, g_n, h) derived from a domain-separated seed. For randomness $r \leftarrow \mathbb{F}$, the commitment is

$$C = \prod_{i=1}^n g_i^{m_i} \cdot h^r \in \mathbb{G}.$$

Under discrete-logarithm hardness in \mathbb{G} , these commitments are computationally binding; they are also perfectly hiding. To avoid linkability, the wallet re-randomizes r across sessions. If it precomputes several reusable commitments, we index them $(C^{(j)}, r^{(j)})$; both offline and online proofs in a session reference the same $C^{(j)}$, allowing the verifier to link the phases without learning \mathbf{m} .

Credential use is captured by two relations:

- *Prepare (offline)*. Once per credential, the wallet verifies σ_I under PK_I , parses S into \mathbf{m} , computes digests $\{h_i\}$, derives a commitment $C^{(j)}$, and produces a reusable proof

$$\pi_{\text{prep}}^{(j)} : “S \text{ parses to } \mathbf{m}, \sigma_I \text{ verifies, and } C^{(j)} \text{ commits to } \mathbf{m}”.$$

- *Show (online)*. For each presentation, the verifier sends a challenge ch . The device signs it as $\sigma_{ch} = \text{Sign}_{SK_D}(ch)$. The wallet proves that all predicates in the verifier’s policy hold with respect to $C^{(j)}$ and incorporates σ_{ch} :

$$\pi_{\text{show}}^{(j)} : “\text{policy holds for } C^{(j)}, \text{ and the session is bound via } \sigma_{ch}”.$$

The verifier checks $\pi_{\text{prep}}^{(j)}$, $\pi_{\text{show}}^{(j)}$, their consistency on $C^{(j)}$, and verifies σ_{ch} under PK_D .

This split amortizes heavy work (signature verification, parsing, commitment) offline, leaving online interaction to short proofs plus one device signature.

Predicates and Policies A *predicate* is a Boolean function $f(\mathbf{m}[S]) \in \{0,1\}$ over a subvector indexed by $S \subseteq [n]$. Typical predicates include range checks ($m_i \geq 18$), equality or membership tests (e.g., m_i equals a country code), and cross-credential comparisons. A *policy* is a finite set of predicates chosen by the verifier. In each session, the wallet proves in zero knowledge that all predicates in the policy hold with respect to $C^{(j)}$, revealing only what the policy requires. Because predicates are modular, the proving backend can be swapped (e.g., from a SNARK to a post-quantum argument system) without changes to issuer infrastructure.

We assume the issuer is honest and operates standard PKI. Verifiers are semi-honest: they check proofs correctly but may collude to compare transcripts. Unlinkability relies on re-randomization of commitments, so the only stable value within a session is $C^{(j)}$, intentionally shared between *Prepare* and *Show*.

3 Our zkID

At a high-level, we propose a generic zkSNARK wrapper over an EUDI digital credential, which will either be issued in SD-JWT format or the mDL data format specified in standard [ISO/IEC 18013-5](#)). The backend proving system we use will be a combination of Spartan and Hyrax commitments, with modifications to be zero-knowledge.

There are two (2) key ideas to highlight within our proposed architecture:

- **Pre-processing batches of re-randomized proofs** of issuer-signature and credential-parsing, to re-use across each new presentation. We call this the **prepare** relation.
- **Committing to the credential disclosures with Hyrax commitments** [WTas⁺17], which allows us to re-use (or “link”) witnesses across circuits “for free”

For the first item, we note that pre-processing proofs for the **prepare** relation is possible because the relation is independent of the presentation, including the choice of disclosures or predicate proofs. Further optimizations can likely be made to only parse the disclosures of certain attributes within the credential if it is known that the Wallet User will rarely or never present certain disclosures to external verifiers.

The second item differs from the linking circuits approach that Google uses [Fas24]. Note that Google verifiably computes a hiding and binding MAC of the shared witnesses as a public output of the circuit, which the verifier checks consistency of in plain. Although this is only a few linear relations, it requires the prover to also commit to their portion of the key to prevent forging. We instead simply manually separate out the disclosures m_1, \dots, m_n into a designated column when committing to the witness, which is already needed to prove the circuit relation. Then, the verifier simply checks consistency of these commitments when verifying each circuit’s proof.

We note that by using Hyrax commitments, our PoC is not post-quantum secure. In particular, post-quantum computers break the discrete-log assumption, which breaks the binding property of the polynomial commitment scheme. Thus, a malicious Prover could potentially make false proofs about their identity. In future work, we hope to incorporate modified Ajtai lattice-based commitments to ensure post-quantum security [HSS24].

Throughout the remainder of this section, we refer to the EUDI’s Wallet User as the “Prover”, the Relying Party as the “Verifier”, and the EUDI Attestation Authority (EAA) as the “Issuer”. Below, we briefly detail a high-level flow of the interaction between the Issuer, Prover, and Verifier.

3.1 Underlying ZK Circuit

In this section, we describe our high-level ZK circuit C underlying the knowledge the prover needs to present to the verifier. Throughout, we refer to the Issuer with variable I , Prover with variable P , and Verifier with variable V (e.g. in subscripts).

We will detail the ZK wrapper around the SD-JWT credential as an example, but the protocol is analogous for other credential formats. Throughout, we will refer to digests as “message hashes” or just “hashes”, and disclosures as “messages”.

We define a circuit C for proving ownership of an anonymous credential. We let our witness $w = S$ be the SD-JWT credential consisting of messages $\{m_i\}_{i=1}^N$, hash salts $\{s_i\}_{i=1}^N$, hashes $\{h_i\}_{i=1}^N$, and an Issuer signature $\sigma_I = \sigma(h_1, \dots, h_N; SK_I)$. Without loss of generality, we assume that the Prover’s public key PK_P is contained in message m_1 of the credential and indexable as $m_1[1]$. We let our instance $x = (PK_I, \{f_i\}_{i=1}^K, \{p_i\}_{i=1}^K)$ contain the Issuer’s public key PK_I , functions f_i over the messages, output predicates p_i , and finally the nonce signature σ_{nonce} for proving device-binding. The f_i can be arbitrary statements we wish to prove about the messages. For example, one could define a function $f_i(m_1, \dots, m_N) = m_1$ would output a predicate that is just the disclosure of message m_1 .

Underlying ZK Circuit C for Verifiable Credential

We define circuit $C(x = (PK_I, \{f_i\}_{i=1}^K, \{p_i\}_{i=1}^K), w_C = (S))$ as follows:

1. Assert $\text{parse}_{\text{SD-JWT}}(S) = (\{m_i\}, \{s_i\}, \{h_i\}_i, \sigma_I)$ parsing of the SD-JWT into messages $\{m_i\}_{i=1}^N$, message salts $\{s_i\}_{i=1}^N$, hashes $\{h_i\}_{i=1}^N$ and Issuer signature σ_I .
2. Assert $h_i = \text{SHA256}(m_i, s_i) \quad \forall i \in [n]$, i.e. that messages hashes correspond to messages and salts
3. Assert $p_i = f_i(m_1, \dots, m_n) \quad \forall i \in [n]$, i.e. correct evaluation of the predicates
4. Assert $\text{ECDSA.verify}(\sigma_I, PK_I) = 1$, i.e. the credential signature verifies under the Issuer public key
5. Assert $\text{ECDSA.verify}(\sigma_{\text{nonce}}, m_1[1]) = 1$, i.e. that the live nonce signature corresponds to the public key the credential was issued to

3.2 Pre-processing and linking proofs

The main speedups from our proving system will come from splitting our high-level circuit C above into two (2) circuits for different relations regarding the digital credential, namely a **prepare** and a **show** relation, analogous to Microsoft’s Crescent Credentials [PPZ24]. This is advantageous because proofs of the **prepare** relation can be computed a-priori for any credential, as they do not depend on the claim being proved at presentation time. Pre-computing these proofs will save significant time per presentation, and reduce the performance bottleneck to that of proving the **show** relation.

One issue that arises is the need to ensure consistency of witnesses across these separate circuits, or what we call “linking proofs”. At a high level, as opposed to Google’s MAC approach [Fas24], the prover sends Hyrax commitments to the parts of the witness re-used across circuits, which ends up being just the raw messages $\{m_i\}_i$. The verifier can then check consistency of these witnesses across the circuits C_i by comparing the Hyrax commitments they receive as part of the proof. This approach gives us linking “for free”, as the Prover already needs to compute these Hyrax commitments as part of the proof.

We highlight that we are no longer splitting up circuits by their field operations (e.g. SHA256 attestations over a binary extension field and an ECDSA verifications over a

prime field), but rather by pre-processing and per-presentation relations. In particular, the circuit for **prepare** will necessarily involve wrong-field arithmetic by including both the SHA256 hashes and the Issuer ECDSA signature verification. However, because of the ability to pre-compute proofs of the **prepare** relation, the more important thing becomes to choose curves that allow for i) efficient show relations and ii) linking the prepare and show relation. Since the verifier can only check equality of Hyrax Pedersen commitments defined over the same curve, we must use the same curve for proving both the **prepare** and **show** relations. Thus we choose a curve with a scalar field equivalent to the base field of the nonce signature σ_{nonce} for efficient signature verification. Because most Hardware Security Modules (HSMs) sign over the P256 curve, we choose the Tom256 (T256) curve for our backend, which has scalar field equivalent to the base field of P256.

We now detail each of the two (2) relations/circuits below.

3.2.1 The prepare relation:

The **prepare** relation checks the validity of issuer signature, parses the SD-JWT, and verifies all the message hashes, none of which depend on the specific presentation. Thus, the prover will periodically pre-compute and store a batch of re-randomized proofs of the prepare relation. These proofs will utilize Hyrax Pedersen vector commitments as introduced above in order to link the proofs of **prepare** relation to the **show**.

Circuit C_1 for the prepare relation

We define circuit $C(x = (PK_I), w_i = S, w = (\{m_i\}_{i=1}^N))$ as follows:

1. Assert $\text{parse}_{\text{SD-JWT}}(S) = (\{m_i\}, \{s_i\}, \{h_i\}_i, \sigma_I)$ parsing of the SD-JWT into messages $\{m_i\}_{i=1}^N$, message salts $\{s_i\}_{i=1}^N$, hashes $\{h_i\}_{i=1}^N$ and Issuer signature σ_I .
2. Assert $h_i = \text{SHA256}(m_i, s_i) \quad \forall i \in [n]$, i.e. that messages hashes correspond to messages and salts
3. Assert $\text{ECDSA.verify}(\sigma_I, PK_I) = 1$, i.e. the credential signature verifies under the Issuer public key

The backend proving system we will use for verifiably computing circuits is Spartan, coupled with a Hyrax-style Pedersen commitment scheme. We can express the circuit computation as some R1CS relation

$$(A \cdot Z) \circ (B \cdot Z) = (C \cdot Z),$$

where $\vec{Z} = (io, 1, \vec{w})$ and io are the public input/outputs. Spartan proves knowledge of a vector Z of length $n := |Z|$ that satisfies the R1CS instance.

To produce zkSNARK proofs for this circuit C_1 , the prover will proceed in two phases:

1. **prepareCommit**: Separates out a column containing only message hashes $\{m_i\}_{i \in [N]}$ in Z and computes an initial Hyrax commitment $c^{(1)} = \{c_i^{(1)}\}_{i \in [\sqrt{n}]}$, which includes a Pedersen commitment to the messages column $c_1^{(1)} = \text{com}(m_1, \dots, m_N; r_1^{(1)}) = g_1^{m_1} \dots g_N^{m_N} g_{N+1}^{r_1^{(1)}}$ with initial randomness $r_1^{(1)}$.
2. **prepareBatch**:
 - (a) Re-randomizes this initial Hyrax commitment to get a batch of commitments $c^{(j)} = \{c_i^{(j)}\}_{i \in [\sqrt{n}]}$, each of which contains a Pedersen commitment to the messages $c_1^{(j)} = \text{com}_1^{(1)} \cdot g_{N+1}^{r_1^{(j)} - r_1^{(1)}}$ for all $j \in [m]$, where our batch size m depends on the frequency of proof generation and demand for the credential

- (b) Continues the Spartan sumcheck IOP on each $c^{(j)}$ to produce a batch of proofs $\{\pi_{\text{prepare}}^{(j)}\}$ for $j \in [m]$ of the **prepare** relation.

The prover will run **prepareBatch** periodically to both generate re-randomized commitments $c^{(j)}$ and store the randomness for the message column commitment $r_1^{(j)}$ for linking purposes, as well as generate and store batches of issuer-signature proofs $\pi_{\text{prepare}}^{(j)}$ that can be used for each presentation.

3.2.2 The show relation:

At a high-level, our show relation will i) verifiably compute any functions f_i over the SD-JWT messages (such as disclosures, range checks, etc.), and ii) check that the credential belongs to the prover's device (also known as proof of "device-binding"). As part of device-binding, the prover will sign a verifier **nonce** outside of the circuit, as outlined in flow ?? . Let us denote this signature by $\sigma_P = \sigma(\text{nonce}; SK_P)$

Again, we will use T256 curve for our backend proving system so that the holder in-circuit signature verification can proceed naturally in the right field.

Circuit C_2 for the show relation

We define circuit $C_2(x = (\{f_i\}_{i=1}^K, \{p_i\}_{i=1}^K), w = \{m_i\}_{i=1}^N)$ as follows:

1. Assert $p_i = f_i(m_1, \dots, m_n) \quad \forall i \in [n]$, i.e. correct evaluation of the predicates
2. Assert $\text{ECDSA.verify}(\sigma_{\text{nonce}}, m_1[1]) = 1$, i.e. that the live nonce signature corresponds to the public key the credential was issued to

As part of computing proof $\pi_{\text{show}}^{(j)}$ for presentation $j \in [m]$, the Prover will once again separate out the messages into a separate column to compute a Hyrax commitment over the Tom256 curve. In particular, the Prover uses *the same* randomness $r_1^{(j)}$ used during the **prepareBatch** process to compute the Pedersen commitment to the messages column. The verifier will then check that the Pedersen commitment to the messages column for $\pi_{\text{show}}^{(j)}$ equals that of proof $\pi_{\text{prepare}}^{(j)}$ for circuit C_1 .

3.3 Adding ZK to Spartan

Our construction uses Circom in the frontend to compile our computation into an R1CS (instance, witness) pair $(x = (\mathbb{F}, A, B, C, io, n, m), \vec{w})$, which we then feed into the Spartan IOP coupled with Hyrax-style Pedersen polynomial commitments.

Recall that our R1CS constraint looks like the following:

$$(A \cdot \vec{Z}) \circ (B \cdot \vec{Z}) - (C \cdot \vec{Z}) = 0$$

where our square matrices A, B, C have size n and $\vec{Z} = (\vec{w}, 1, io)$.

Recall that Spartan converts an R1CS constraint into the following zero-check:

$$\sum_{x \in \{0,1\}^{\log n}} \tilde{e}q(x, \tau) \left[\left(\sum_{y \in \{0,1\}^{\log n}} \tilde{A}(x, y) \tilde{Z}(y) \right) \left(\sum_{y \in \{0,1\}^{\log n}} \tilde{B}(x, y) \tilde{Z}(y) \right) - \left(\sum_{y \in \{0,1\}^{\log n}} \tilde{C}(x, y) \tilde{Z}(y) \right) \right] = 0$$

for some random challenge $\tau \in \mathbb{F}$

There are two components in Spartan that we need to modify to be ZK. The first is making the sumchecks ZK. The second is to ensure that the opening \tilde{Z} using the commitment to \vec{Z} does not leak information about our witness \vec{w} .

3.3.1 Adding ZK to sumcheck

The Spartan protocol consists of several sumchecks in parallel and operates over some field \mathbb{F} . There are various existing techniques to make sumcheck ZK. We employ one using similar methods as in Zhang et. al. [ZXZS19], which adds uniformly random pads to the sumcheck transcript.

In particular, suppose at each round i of the sumcheck protocol, the prover sends over $s_i(X) := \sum F(r_1, \dots, r_{i-1}, X, x_{i+1}, \dots, x_m)$ where r_i is the Verifier challenge sent for round i . Then instead of having the verifier check the sumcheck, the prover will prove in ZK that the unpadded transcript satisfies the verifier's (linear) checks. To do this, the prover will need to commit to the uniformly random pads ahead of time. Then, as long as the Fiat-shamir challenges is generated from the transcript including these random pad commitments, the prover cannot simply lie about the pads to satisfy the sumcheck relation.

Adding ZK to sumcheck

1. Prover commits to pads $R_i(X) \xleftarrow{\$} \mathbb{F}_1[x]$ for all $i \in [\log n]$. These are linear polynomials, and can thus be represented by its two coefficients $R_i[0]$ and $R_i[1]$.
2. Instead of sending partial sums

$$s_i(X) := \sum_{(x_{i+1}, \dots, x_n) \in \{0,1\}^{n-i}} F(r_1, \dots, r_{i-1}, X, x_{i+1}, \dots, x_n)$$

for each round of sumcheck, the Prover sends polys $s'_i(X) = s_i(X) + R_i(X)$, essentially a one-time-padded transcript.

3. It suffices to show the following linear relation in zero-knowledge,

$$\left[\begin{array}{c|c|c} A & B & C \end{array} \right] \begin{bmatrix} \vec{S} \\ \vec{R} \\ C \\ F(r) \end{bmatrix} = \vec{0}$$

where

$$\vec{S} = [s'_1[0], s'_1[1], \dots, s'_n[0], s'_n[1]]^\top,$$

is the column vector of sumcheck transcripts such that $s'_i = s'_i(X) = s'_i[0]X + s'_i[1]$, and

$$\vec{R} = [R_1[0], R_1[1], \dots, R_n[0], R_n[1]]^\top,$$

is the column vector of random pads, and where matrices A, B, C are given by

$$A = \begin{bmatrix} 1 & 2 & 0 & 0 & 0 & 0 & 0 & \cdots & 0 & 0 & 0 & 0 \\ -r_1 & -1 & 1 & 2 & 0 & 0 & 0 & \cdots & 0 & 0 & 0 & 0 \\ 0 & 0 & -r_2 & -1 & 1 & 2 & 0 & \cdots & 0 & 0 & 0 & 0 \\ \vdots & \vdots & \vdots & \vdots & \vdots & \vdots & \vdots & \ddots & \vdots & \vdots & \vdots & \vdots \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & \cdots & -r_{n-1} & -1 & 1 & 2 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & \cdots & 0 & 0 & r_n & 1 \end{bmatrix}$$

$$B = \begin{bmatrix} -1 & -2 & 0 & 0 & 0 & 0 & 0 & \cdots & 0 & 0 & 0 & 0 \\ r_1 & 1 & -1 & -2 & 0 & 0 & 0 & \cdots & 0 & 0 & 0 & 0 \\ 0 & 0 & r_2 & 1 & -1 & -2 & 0 & \cdots & 0 & 0 & 0 & 0 \\ \vdots & \vdots & \vdots & \vdots & \vdots & \vdots & \vdots & \ddots & \vdots & \vdots & \vdots & \vdots \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & \cdots & r_{n-1} & 1 & -1 & -2 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & \cdots & 0 & 0 & -r_n & -1 \end{bmatrix}$$

$$C = \begin{bmatrix} -1 & 0 \\ 0 & 0 \\ 0 & 0 \\ \vdots & \vdots \\ 0 & 0 \\ 0 & -1 \end{bmatrix}$$

4. The Prover computes a public random challenge α (e.g. hashing the transcript) and compresses the relation into a dot product

$$[\vec{u}] \begin{bmatrix} \frac{\vec{S}}{\vec{R}} \\ \frac{C}{F(r)} \end{bmatrix} = \vec{0} \quad (1)$$

where $\vec{u} = [1, \alpha, \alpha^2, \dots, \alpha^n]$ $\begin{bmatrix} A & B & C \end{bmatrix}$ is a random linear combination of the rows.

5. The Prover and Verifier engage in a proof-of-dot product protocol to prove the relation above, such as an Inner Product Argument used in Bulletproofs [BBB⁺17]

3.3.2 Adding ZK to the opening of \tilde{Z}

In order to add ZK to the opening of, we simply “append” Z with uniformly random pads. Specifically, we assign random evaluations $Z(x) \xleftarrow{\$} \mathbb{F}$ on any remaining point in the hypercube $x \in \{0, 1\}^{\log n}$. Then, we see that $\tilde{Z}'(x_1, \dots, x_{\log n}) = \tilde{Z}(x_1, \dots, x_{\log n}) + eq(r, x)Z(x)$. Note that now $\tilde{Z}'(x_1, \dots, x_{\log n})$ is distributed uniformly at random. If $n = |Z|$ is not already a power of 2, then we can simply fill at least one of the remaining evaluations on $\{0, 1\}^{\log n}$ with a single random pad. If $n = 2^m$, we can add another dimension to the hypercube of evaluations of Z .

3.4 Cost analysis

The following section computes the Prover and Verifier costs of Spartan instantiated with Hyrax Pedersen commitments on R1CS instances $(x = (\mathbb{F}, A, B, C, io, n, m), \vec{w})$, where io denotes the vector of public inputs/outputs, $n = |\vec{w}| + io + 1$ is the dimension of our matrices, and m is the number of nonzero entries in our matrices A, B, C . We let $Z = (\vec{w}, io, 1)$. It is often reasonable to assume that our R1CS matrices are sparse, i.e. $m = O(n)$. However, we present the costs below independent of this assumption.

- **Prover time:** (1) $O(m)$ to generate sumcheck transcript, (2) $O(m)$ to evaluate MLEs of A, B, C , (3) $O(n)$ to commit to the MLE of Z (computing \sqrt{n} MSMs of size \sqrt{n}) and opening the MLE of Z , for a total cost of $O(m)$.

- **Proof length:** (1) $O(\log n) \cdot |\mathbb{F}|$ length of the sumcheck transcript, (2) $O(\sqrt{n}) \cdot |\mathbb{G}|$ length commitment to MLE of Z , (3) $O(\log n) \cdot |\mathbb{G}|$ length of argument opening MLE of Z , for a total length of $O(\sqrt{n})$ group or field elements.
- **Verifier time:** (1) $O(\log n)$ to verify sumcheck transcript, (2) $O(m)$ to evaluate the MLEs of A, B, C (with sparse commitment scheme and memory checking), (3) $O(\sqrt{n})$ to open the MLE of Z , for a total of $O(m + \sqrt{n})$.

With the ZK modifications to Spartan, we can see the asymptotic costs remain the same, as follows:

- **Prover time:** additionally computes $O(\log n)$ constant-size commitments to the sumcheck transcript pads r_i , and $O(\log n)$ engages in new sumcheck relation IPA (or some other ZK dot product argument) for vector of length $O(\log n)$, which still gives $O(m)$ prover work.
- **Proof length:** sumcheck and openings are the same length but just padded, but added on $O(\log n)$ size $|\mathbb{G}|$ commitments, and a length $O(\log \log n)$ sumcheck relation IPA proof, which still gives a proof length of $O(\sqrt{n})$ group or field elements.
- **Verifier time:** no longer needs to do $O(\log n)$ (sumcheck), but still needs $O(m)$ (evaluating MLEs of A, B, C) + $O(\sqrt{n})$ (opening MLE of Z) + $O(\log n)$ for sumcheck relation IPA verification, which still gives $O(m + \sqrt{n})$ runtime.

3.5 Security analysis

The correctness follows immediately from the correctness of the Spartan SNARK and the fact that the Prover uses the same randomness for the Hyrax commitments across the **show** and **prepare** circuits for each presentation i .

The soundness of our protocol follows from the soundness of Spartan. In particular, we can extract the full witness credential from the **prepare** relation.

Intuitively, zero-knowledge follows from the hiding property of the commitment scheme as well as the zero knowledge property of the Spartan zkSNARK proving system; For proof i , simulator can randomly sample the linked commitment com_i both distributions to reuse across both proofs, both in the commitment itself and also in the IPA used to open the Hyrax commitment to $Z(r_1, \dots, r_{\log n})$. We can show that this commitment is independent of the rest of the view of the Verifier, which consists of the following:

- Sumcheck polynomials $\{s'_i(X)\}_{i \in [\log n]}$ for each of the sumchecks in Spartan
- $\{r_i\}$ Fiat-Shamir challenges during the sumcheck
- Transcript from the IPA on the sumcheck relation in ZK
- $\{com(z_i)\}_{i \in [\sqrt{n}]}$ Hyrax commitment to Z , which involves a Pedersen commitment to each of the columns of a $\sqrt{n} \times \sqrt{n}$ matrix representation of \vec{Z}
- Transcript from the IPA for opening $Z(r_1, \dots, r_{\log n})$
- The claimed value of $Z(r_1, \dots, r_{\log n})$

Since we appended random pads to \vec{Z} in our ZK modification in Section 3.3.2, the distribution of $Z(r_1, \dots, r_{\log n})$ is random and independent of Z , and therefore independent of $\{m_i\}_i$. Furthermore, $s'_i(X)$ have totally random pads on them and their distribution is independent of Z , and therefore independent of $\{m_i\}_i$. Assuming the hiding property of the Pedersen commitment schemes for messages sent during an IPA, we can also use

the simulators for the IPAs without changing their joint distribution with the rest of the transcript.

Then, we can simply run the piece-wise simulators for each zkSNARK proof for circuits C_1 and C_2 to simulate the remainder of the view.

4 Experiments

5 Application to EUDI

Within the EUDI Architecture and Reference Framework [Eur23], the practical question is how to introduce zero-knowledge capabilities without disrupting established roles, formats, and certification paths. This section states how the construction fits that setting and what trade-offs it entails. Subsections are structured according to [Topic G](#) in the EUDI ARF discussion thread.

Throughout this discussion, we continue to refer to EUDI’s Wallet User as the “Prover”, the Relying Party as the “Verifier”, and the EUDI Attestation Authority (EAA) that acts as a PID/Attestation Prover as the “Issuer”.

Issuance. The construction is designed to wrap existing credential encodings rather than replace them. It accommodates SD-JWT and ISO/IEC 18013-5 mDL so that wallets and relying parties retain current disclosure grammars and parsing logic. Issuers remain oblivious to the use of zkSNARKs; no changes to issuance pipelines or device secure elements are required, and Issuers also maintain exclusive control of their private keys. The proof layer is circuit-defined and therefore highly programmable, which allows our scheme to easily adapt to future Issuer-side migrations (for example, a change of signature scheme) by simply updating the Prover circuit and public parameters, rather than introducing new format-specific protocols. The approach interoperates with current public-key infrastructure based on ECDSA or RSA and does not prescribe a switch of algorithm or hardware.

Efficiency. Proving is split into two relations. A fixed relation captures Issuer-signature verification, credential parsing, and commitment preparation; it runs infrequently and is amortized per credential. A live, presentation-specific relation captures the disclosures and predicates for a single session; it runs per presentation. This separation aims to keep Prover time and memory costs within typical web and mobile budgets, and to bound latency where most critical: at the time of Prover-Verifier interaction. The proof system and commitment layer are modular, so improvements in either component can be adopted without redesigning the higher-level flow. In contrast to other designs, proofs of Issuer-signature verification and parsing are pre-computed offline to reduce work during live Prover-Verifier interaction.

Discussion. The present instantiation follows the Spartan line and relies on sumcheck and Hyrax-style Pedersen commitments under the Discrete Log assumption, rather than pairing-based assumptions; there is no universal trusted setup. We avoid pairing-friendly curves and the operational burden of a trusted setup ceremony across many Provers, Issuers, Verifiers, and other independent bodies. The construction is not currently post-quantum secure, but the modular structure leaves a path to replacing the commitment scheme layer with lattice-based alternatives as they mature. Some components have not yet been standardized; we note this is a shared condition across competing approaches that we call out explicitly. Regarding standardizations: sumcheck is a highly well-known protocol with information-theoretic security independent of cryptographic assumptions;

Pedersen commitments have been used since 1991 [Ped92] and rely only on the discrete-log assumption, which standardized ECDSA signatures already rely on.

Summary for EUDI. The design aligns with Annex 2 format expectations (see Section 8), requires no changes to Issuers, supports current PKI deployments, and separates fixed from presentation-specific work to keep live presentation costs low. It avoids pairing-based assumptions and a universal setup, and leaves a path to future cryptographic upgrades without disrupting wallet or Issuer operations.

6 Security

In our security model, we assume that the Prover is malicious, and that each Verifier is semi-honest, meaning that if the Prover presents a valid proof that they own a credential with some property, the Verifier will grant access to any services for which the property suffices.

Verifier’s side For security on the Verifier’s side, our soundness analysis considers the probability that a malicious Prover without real ownership of a valid credential can generate a false proof of ownership.

Prover’s side For security on the Prover’s side, we guarantee that our proofs are zero-knowledge, so that a semi-honest and computationally-bounded Verifier cannot get any additional information about the Prover’s credential beyond what is publically revealed in the proof. In particular, we do not consider the case where the Verifier is malicious during presentation, e.g. where a false Verifier pretends to be an authorized Verifier. The problem of Verifier identity lies outside the scope of this paper.

Furthermore, we assume that Verifiers can collude with each other, i.e. that Verifiers V_1, \dots, V_N that have received proofs $\{\pi_1\}, \dots, \{\pi_N\}$ from a given Prover P can compute functions $f(\pi_1, \dots, \pi_N)$. Therefore, we desire the **unlinkability property**: given π_1, \dots, π_N , the Verifiers should not be able to determine whether or not any two of these proofs came from the same Prover P . Note that this requires the Prover to re-randomize each presentation’s proof; a static zero-knowledge proof of the same statement, while not revealing private credential information, would still look the same across presentations. In that case, it may be possible for the Verifier to de-anonymize a Prover by linking their “anonymous” activity across presentations and analyzing metadata, e.g. time of presentation. Fortunately, our scheme is unlinkable due to the re-randomization of proofs between each presentation. By the zero-knowledge property for each presentation, we can simulate the distribution of proofs without knowledge of the witness. To simulate an entire set of proofs received by distinct colluding Verifiers, we can independently simulate each proof.

Finally, as our scheme is currently presented in Section 3, we assume that Verifiers will not collude with Issuers even though they can see the Issuer public key. To bypass this assumption and prevent Issuer tracking in the case of malicious Verifiers that collude with Issuers, we propose here the maintenance of a trusted Merkle tree on trusted Issuer public keys. Then our Prover’s circuit would prove knowledge of a valid Issuer signature from some key in the Merkle tree, and the public input/output would just be the Merkle root rather than any specific Issuer public key.

Both Prover and Verifier security When modelling the Issuer, we assume that the Issuer is trusted during issuance by both the Prover and Verifier, i.e. will not Issue false credentials or sell personal information that is necessarily to obtain about individuals to issue a credential.

6.1 Other considerations

Our scheme currently does not require any interaction from the Issuer for credential presentation beyond initial issuance.

However, our scheme does require the use of internet access (without, there will be risks with authorizing someone before their credential can be checked against the current state). In the case (as presented) where Issuer public key is a public input/output, we assume there is an online registry of trusted Issuer keys that the Verifier can check the proof against. This requires live internet access in the same way that credit card transactions do, in order to check the most current registry of public keys. Even in the case of a Merkle inclusion proof, where the Issuer key is also private, the Verifier would need to check that the public Merkle root matches the trusted root stored online. It is possible to store encrypted transactions/credential presentations to be checked later once internet access is restored, in the same way as offline credit card transactions do. However, there are necessary risks with this approach; it would be up to the specific vendor and/or service prover what levels of risk can be tolerated from delayed credential authentication. For example, some service provider (Verifier) may be fine only periodically downloading the current registry of trusted Issuer keys (and/or Merkle roots) and simply checking against their last downloaded version before granting access.

Finally, as mentioned previously, our scheme is not quantum resistant due to the use of Hyrax commitments. Again, we believe this is easily fixable with the introduction of modified Ajtai lattice-based commitments, which are post-quantum secure.

7 Conclusion

References

- [BBB⁺17] Benedikt Bünz, Jonathan Bootle, Dan Boneh, Andrew Poelstra, Pieter Wuille, and Greg Maxwell. Bulletproofs: Short proofs for confidential transactions and more. Cryptology ePrint Archive, Paper 2017/1066, 2017.
- [Eur23] European Commission. The european digital identity wallet architecture and reference framework. Technical report, European Commission, 2023.
- [Fas24] Matteo Frigo and abhi shelat. Anonymous credentials from ECDSA. Cryptology ePrint Archive, Paper 2024/2010, 2024.
- [fS21] International Organization for Standardization. Iso/iec 18013-5:2021 personal identification — iso-compliant driving licence part 5: Mobile driving licence (mdl) application, 09 2021.
- [FYC25] Daniel Fett, Kristina Yasuda, and Brian Campbell. Selective disclosure for JWTs (SD-JWT). Technical Report draft-ietf-oauth-selective-disclosure-jwt-22, IETF OAuth WG, May 2025. Internet-Draft.
- [GRS⁺] Paul A. Grassi, Justin P. Richer, Sarah K. Squire, James L. Fenton, and Ellen M. Nadeau. Digital identity guidelines: Federation and assertions (nist special publication 800-63c). NIST Special Publication 800-63C SP 800-63C, National Institute of Standards and Technology (NIST). Includes updates as of 2020-03-02; Privacy Authors: Naomi B. Lefkowitz, Jamie M. Danker; Usability Authors: Yee-Yin Choong, Kristen K. Greene, Mary F. Theofanos.
- [HSS24] Intak Hwang, Jinyeong Seo, and Yongsoo Song. Concretely efficient lattice-based polynomial commitment from standard assumptions. Cryptology ePrint Archive, Paper 2024/306, 2024.

- [Ped92] Torben P. Pedersen. Non-interactive and information-theoretic secure verifiable secret sharing. In Joan Feigenbaum, editor, *Advances in Cryptology – CRYPTO’91*, volume 576 of *Lecture Notes in Computer Science*, pages 129–140, Santa Barbara, CA, USA, August 11–15, 1992.
- [PPZ24] Christian Paquin, Guru-Vamsi Policharla, and Greg Zaverucha. Crescent: Stronger privacy for existing credentials. Cryptology ePrint Archive, Paper 2024/2013, 2024.
- [Wor25] World Wide Web Consortium (W3C). Verifiable credentials data model v2.0. W3C Recommendation REC-vc-data-model-2.0, World Wide Web Consortium (W3C), W3C, May 2025. Former editors: Grant Noble, Dave Longley, Daniel C. Burnett, Brent Zundel, Kyle Den Hartog.
- [WTas⁺17] Riad S. Wahby, Ioanna Tzialla, abhi shelat, Justin Thaler, and Michael Walfish. Doubly-efficient zkSNARKs without trusted setup. Cryptology ePrint Archive, Paper 2017/1132, 2017.
- [ZXZS19] Jiaheng Zhang, Tiancheng Xie, Yupeng Zhang, and Dawn Song. Transparent polynomial delegation and its applications to zero knowledge proof. Cryptology ePrint Archive, Paper 2019/1482, 2019.

8 Appendix: EUDI Annex 2 Requirements

This section is devoted to a review of the EUDI ARF’s Annex 2, which covers high-level requirements for the EUDI Wallet. The full Annex can be found [here](#).

In this Appendix, we address all of the Topics presented in Annex 2 of the EUDI Architecture Reference Framework. The first part of this section is addressed to a general audience; it translates the vocabulary used in the EUDI ARF to cryptographic terms that those reading this paper might be familiar with.

We then present some additional directions for our proof-of-concept that introduces potential solutions to some of the Topics (requirements) in the EUDI ARF. Finally, we lay out an exhaustive table addressing each of the Topics of the ARF, indicating: (i) whether or not our paper addresses and satisfies the specifications (see “Within Scope” column), (ii) any potential privacy concerns that may arise from satisfying the specification of the Topic (iii) a brief high-level proposals for any privacy concerns that do arise.

The following table consists of Topics that we either directly tackle in our proof-of-concept, or are actively working on.

Topic	Summary	Comments	Within Scope	Paper Notes	Privacy Concerns?	Privacy Notes
1	Device binding and remote flows, according to [OpenID4VP] standard	Some of the points (e.g. verifying relying party identity, UX flows for which credential to present) are not strictly covered within our flow	Yes	Device binding: we propose that WSCA signature of nonce should be checked against pk of credential PRIVATELY in-circuit (otherwise reveals pk → linkability). Relying party verifying PID/QEAA signatures from trust list: can check issuer pk from list while hiding issuer signature check PRIVATELY in-circuit	Yes	We would propose merkle inclusion proof to hide the specific issuer pk (in case that issuer has only issued a few credentials), but would rely on external trusted maintenance/agreement of the merkle tree. This may not strictly align with "validate signature using trust list".
2	Wallet must support mDLs		Yes	Must support mDLs (specified in [ISO/IEC 18013-5]). Crescent already supports mDLs CBOR parsing	Maybe	
7	Only issuers can revoke, using an "attestation status/revocation list mechanism", that relying parties also use	Either short term credentials, have an attestation status (e.g. suspension) list mechanism, or attestation revocation list mechanism	Ongoing work	Prover needs to provide some kind of proof of non-revocation. We propose providing merkle inclusion proof of a public online list	Yes	Without attestation lists being public, would either i) need to phone home to issuer to see the status of credential/obtain a proof → issuer surveillance, or ii) provide an ID the relying party can check against a public list → linkability
10	Wallets must support proximity and remote [OpenID4VP] flows. Wallets must support mDLs and SD-JWTs. UX flow around user accepting newly issued PID/attestation		Yes	Again must support mDLs (specified in [ISO/IEC 18013-5]). Crescent already supports mDLs CBOR parsing		
11	Pseudonyms issued by a pseudonym provider. Allows relying party to recognize users across presentations		Yes	To eliminate need for an external pseudonym provider + allow for multiple pseudonyms controlled by the user: we propose computing/outputting a deterministic nullifier hash $H(\text{public_key}, \text{random_salt})$ (where the wallet stores the random_salt) in-circuit, to use as the pseudonym	Yes	If pseudonym provider is issuer, this is really bad (can track full identity whenever issuer-assigned pseudonym used). If pseudonym provider is an external party, perhaps need some kind of id disclosure to get a pseudonym → similar to issuer tracking

Topic	Summary	Comments	Within Scope	Paper Notes	Privacy Concerns?	Privacy Notes
17	Ensuring users with multiple accounts/credential signins are actually the same person — requesting linking IDs		Yes	Credentials should be re-randomized always	Yes	"Request the identified EUDI Wallet User to identify with another eID means which is accepted by the Relying Party so to link the data received from the EUDI Wallet with the account to which the User proved to have access to" → further supports need for efficient re-randomization
18	User presenting info/proofs across multiple credentials. Shall request proof-of-association of PKs each credential is issued to from the WCSA/WCSD	Here is a potential Schnorr-style ZKP for association between PKs. Main idea: proof of knowledge of dlog relationship btw the two	Yes	Proof-of-association should not have PKs in plaintext to the verifier (otherwise provides linkability). Need wallet user to do process proof received directly from WCSA. We propose an in-circuit verification of the ZKP received from WCSA (as a recursive proof) + matching of PKs to the ones in the credential	Yes	« «
29	Should allow for issuance of eIDs that let someone represent someone else. Not speced out yet, ad-hoc	Unclear	Potential future work	We propose including both entities' PKs in the credential where transaction log would show up in represented person's as well. Maybe phone home to person being represented is OK?	Maybe	« «
35	Protocol for PID issuance	Yes	Yes	See privacy concerns around wallet user providing WTE for PID/attestation issuer to verify before credential issuance	Yes	In the cases that the issuer is just another verifier or EUDI wallet instance (user) — want WTE verification to be private (especially for credential issuances that don't require sensitive info/verification) for unlinkability

Topic	Summary	Comments	Within Scope	Paper Notes	Privacy Concerns?	Privacy Notes
38	Need to ensure use of Wallet Instance Attestations (WIA) to relying parties does not allow tracking. Wallets must revoke/suspend when PID issuer asks		Yes		Yes	If WIAs are a separate list and just a static watermark, then can do merkle inclusion proof for unlinkability. We recommend against publishing list of public keys as a WIA mechanism — since this causes device-binding to function as doxxing (can check signatures against each of the potential public keys)

For the following Topics in the EUDI ARF, we assume the underlying architecture that the Topic addresses already satisfies the requirements, so that one can add our zkID solution as the credential presentation flow layer on top. We still provide comments, including notes about potential privacy concerns we have about the specification.

This first table presents the Topics that we believe present the most major privacy concerns, and we provide comments on what the threats are.

Topic	Summary	Comments	Privacy Issues	Privacy Notes
9	Wallet provider provides WTE certificate to wallet instances (testifying security of WSCA/D), wallet gives new PK and proof-of-association with WTE key to issuers to receive credentials to	Wallet provider responsibility, WTE is never shown to relying parties (see Topic 18 for proof-of-association sent to relying parties)	Yes	Potential issuer collusion allows for reconstruction of a superset of IDs if WTE is provided in plain. Proof-of-association with keys should also be private/in-circuit if possible
19	Must have overview of all transactions executed through the Wallet Instance that cannot be deleted...		Yes	Need to ensure this is hidden behind some kind of biometric authentication within the wallet, so that cannot be accessed if stolen. Access to this data allows linkage everywhere. Definitely should not be external
33	Backup devices (e.g. HSM backups)		Yes	How do HSM backups work? How to ensure they are secure so that credentials aren't stolen. In the process of backing up, must transfer some sensitive information (e.g. entire credential)

This second table presents the remaining Topics, with comments on why the topic lies within a separate layer from our zkID presentation flow.

Topic	Summary	Comments	Privacy Issues	Privacy Notes
3	PID rulebook	Issuer responsibility	No	
4	mDL rulebook	Issuer responsibility	No	
6	Relying party authentication	User checks relying party ID and certificates, e.g. signature checks		

Topic	Summary	Comments	Privacy Issues	Privacy Notes
12	Should standardize attribute identifiers/syntaxes across namespaces and attestation types for max interoperability	Issuer responsibility / for verifier request convenience. But also somewhat for user legibility/transparency	No	
16	Wallet should allow user to create (qualified electronic) signatures over documents		No	
24	UX flows for proximity (e.g. mDL) - including user approval for disclosure	Just UX layer	No	
25	Similar to topic 12, standardizing vocab for attestation attributes	Issuer responsibility	No	
26	For these standards, anyone can contribute, just be reasonable	Issuer/standards responsibility	No	
27	Relying party should get certificates to be verified by the wallet user before disclosure	Relying party + trusted anchor responsibility	No	
28	Wallet for legal person (e.g. corps, governments, and NGOs) should be diff from natural person (human)	Issuer responsibility	No	
30	Wallet user can also be verifier	Introduces no new requirements for presentation flow	No	
31	Public, no-auth trusted list of certificate issuers (for issuers, wallet providers, and relying parties)	Nothing concerning! Public list is good	No	
42	Qualified trust service providers (qtsp)(issuers of QEAA attestation types) need access to authentic sources (e.g. databases of legal citizens)		No	
43	Policy for which items to be disclosed in certain presentation situations	Needs to be built into the frontend logic before circuit proofs on the user side.	Maybe	Only if some metadata about Wallet (that would provide linkability) can be tracked when requesting info from relying party...
48	Users being able to delete presented data to verifiers		No	Good for privacy!
50	Protocols for reporting relying party abuse	[?] Isn't relying party certificate + checking against trusted anchor list enough?	No	

Finally, the following topics had no high-level requirements (as of July 2025), or their high-level requirements are already contained in other Topics:

Topic	Topic name
5	EUDI Wallet Design Guide
8	Design Solutions on Data Sharing scenarios
13	Developing an EUDI Wallet Architecture Based on Secure Element
14	Developing an EUDI Wallet Architecture Based on External Token
15	Developing an EUDI Wallet Architecture Based on Remote HSM
20	Strong User (Customer) Authentication in the context of electronic payments
21	Diplomas with EUDI Wallet
22	Digital Travel Credentials with EUDI Wallet
23	PID issuance and (Q)EAA issuance
32	PID interoperability
34	Migrate to a different wallet solution

Topic	Topic name
36	Risk Analysis of the EUDI Wallet Usage
37	QES — Remote Signing — Technical Requirements
39	Wallet to wallet technical Topic
40	Reserved
41	Minimum requirements on PuB-EAAs rulebooks
44	QEAA evaluation requirements
45	QEAA Rulebook
46	Protocols and interfaces for Presentation of PID and (Q)EAA with relying parties
47	Protocols and interfaces for PID and (Q)EAA issuance, and (non-)qualified certificates issuance
49	Protocol and interfaces for requesting data deletion to relying parties