zkID technical report

Jane Doe[1,2] and John Doe[1]

[1] Institute A, City, Country, `jane@institute`
[2] Institute B, City, Country, `john@institute`

Abstract. Main deliveries: 1. Technical report on zk component for the digital id wallet 2. A comparison with current works 3. Applying to EUDI.

Keywords: Anonymous credential · programmable zkp

1 Introduction

People use digital credentials to prove a single fact—such as being over 18, holding a license, or belonging to an organization—without presenting an entire document. In the W3C model [Wor25], an issuer signs a credential, a holder keeps it (typically in a wallet application), and a verifier checks a presentation—that is, data derived from one or more credentials and shown to a specific verifier for a specific session. This is the baseline flow we adopt.

A practical privacy risk arises when the same credential is shown to different services over time. If a presentation exposes a stable technical trace—for example, an identifier or handle that repeats—then services can compare logs later and infer that the sessions likely came from the same source. We call this linkability: separate presentations that should remain independent become easy to connect. Standards documents explicitly caution against this: the W3C Verifiable Credentials Data Model states that securing mechanisms must not leak information that would enable a verifier to correlate a holder across multiple presentations [Wor25]; and NIST's federation guidance recommends pairwise pseudonymous identifiers - meaning unique, opaque identifiers for each relying party - to ensure separation across services and reduce the risk of user tracking [GRS+]. Deployments typically provide two recurring capabilities: device binding (confirming that the credential is presented by the authorized wallet) and revocation/status (learning whether the credential is still valid). To support these, systems often expose stable values. For device binding, some ecosystems tie a credential to a long-lived cryptographic key under the control of the device so a verifier can accept only "the right wallet"; if the same key reference or its proof is visible across sessions, different verifiers can match their logs [FYC25]. For status, verifiers query whether a credential was suspended or revoked; a one-to-one status reference (for example, a unique status URL per credential) lets the status provider observe which verifier asked about which holder and when [Wor25]. These two mechanisms keep systems reliable and usable, but because they rely on fixed values, they also leave repeating patterns that make it possible to link separate credential uses.

We preserve the operational stability that deployments need while preventing correlators at verification. Concretely, zkID adds a presentation-layer wrapper—a thin layer at the presentation step (in the W3C sense) that re-randomizes what the verifier observes. After a one-time *prepare* when the wallet stores a credential, each session produces a fresh, per-session randomized proof; the verifier, in addition to its normal signature and policy/status checks, performs one extra consistency check to confirm that the randomized proof matches the underlying credential. We analyze unlinkability under a practical model in which

holders may be malicious and verifiers may collude by comparing what they observe across sessions; issuers do not collude with verifiers, and issuance flows and keys remain unchanged. Under this model, presentations remain verifiable for services while no reusable handle is exposed to link sessions—consistent with the W3C requirement [Wor25].

Our design does not require changes to issuers. Modifications occur only at the two endpoints of presentation: (i) the wallet runs a one-time prepare when storing a credential and, at each use, generates a randomized proof so two verifiers do not observe the same handle; and (ii) the verifier adds one consistency check alongside its existing validation to confirm that the randomized proof aligns with the presented credential. We evaluate zkID under these deployment constraints.

1.1   Related Work

Anonymous Credentials from ECDSA.   Matteo Frigo and Abhi Shelat (Google) propose an anonymous credential scheme for legacy-deployed ECDSA that preserves deployability on existing infrastructures (e.g., P-256, SHA-256) without issuer or device changes and without trusted setup, addressing reluctance to upgrade infrastructures that only support RSA/ECDSA [Fas24]. Their design tackles the bottleneck of zero-knowledge proofs over non-NTT-friendly curves by building a proof stack around sum-check and the Ligero argument [AHIV17], introducing specialized ECDSA circuits and efficient Reed–Solomon encodings, and adding a witness-consistency mechanism to keep common witness values aligned across arithmetic over $\mathbb{F}_p$ (ECDSA) and $GF(2^k)$ (SHA-256). As reported, the system generates a proof for a single ECDSA signature in 60 ms and a zero-knowledge proof for the ISO/IEC 18013-5 mDL presentation flow in 1.2 s on mobile devices [Fas24, §5.3,§6.2]. Device binding follows the mDL live-challenge pattern. The main trade-off is comparatively larger proofs and higher verifier workload than succinct CRS-based SNARKs.

Crescent Credentials.   Christian Paquin, Guru-Vamsi Policharla, and Greg Zaverucha introduce Crescent to upgrade privacy of existing credentials (JWTs, mDL) with selective disclosure and unlinkability without changing issuance processes [PPZ24]. Crescent achieves fast online presentations with compact proofs, while relying on a two-phase workflow:

- Prepare (offline, once per credential): verify issuer signatures, decode and parse the credential into attributes, and produce a Pedersen vector commitment via a Groth16 proof. Reported timings are approximately 27 s for a 2 KB JWT and 140 s for an mDL, with universal-setup parameters of size approximately 661 MB–1.1 GB [PPZ24, §4].

- Show (online, per presentation): re-randomize the Groth16 proof and attach a few proofs over committed attributes; typical latencies are approximately 22 ms for JWT and 41.2 ms for mDL, with proofs around 1 KB. Optional device binding increases these costs (e.g., about 315 ms Show, about 184 ms verify; proof size about 15 KB) [PPZ24, §4].

Crescent exposes a modular committed-attribute interface that admits sub-provers for range checks, credential linking, and binding to session context or device by proving knowledge of a signature under a committed public key. The principal trade-offs are reliance on pairing-based Groth16 with a trusted setup and large universal parameters, and a comparatively heavy (but amortized) offline Prepare phase; in its current form Crescent is not post-quantum secure [Gro16].

zkID.   zkID is a modular, transparent (no-CRS) zero-knowledge wrapper for standardized credentials, prioritizing unlinkability and deployability. We adopt the same two-phase

workflow as Crescent but avoid a universal setup, and we retain ECDSA compatibility as in Google's design while moving issuer verification to the offline phase to reduce online verifier work. Key distinctions from prior work include:

- Transparency: no trusted setup; standard discrete-log assumptions;

- Deployment compatibility: interoperates with mDL/SD-JWT and existing PKI (ECDSA or RSA) without changes to issuer processes or secure elements.

- Avoidance of pairing-based cryptography: unlike pairing-based schemes, zkID avoids pairing-friendly curves.

- Verifier cost: reduces verifier work relative to prior ECDSA-based designs under our constraints while preserving unlinkability under the stated threat model.

In summary, zkID shows that a modular, standards-aligned zero-knowledge wrapper over existing credential formats can be built. Compared to pairing-based schemes, it avoids dependence on pairing-friendly curves; compared to universal-setup SNARK wrappers, it avoids a trusted setup; and relative to prior ECDSA-based designs, it reduces verifier cost while preserving unlinkability under our model. The construction is transparent and PQC-agile; the current instantiation is classical and not quantum-resistant.

In our evaluation, zkID achieves practical performance on commodity hardware. For a typical credential with ..., the one-time *Prepare* phase completes in ... seconds on a standard laptop (or mobile) (..., ...GB RAM), while the per-presentation *Show* phase takes less than ... ms on a mid-range smartphone. Proof sizes are approximately ... KB for the full protocol, and verifier time remains below ... ms on a web server. These results demonstrate that zkID meets latency and resource targets for mobile and web deployments, while preserving unlinkability and avoiding trusted setup. Detailed benchmark results and further breakdowns will be provided in Section 4.

2   Preliminaries – WIP

Notation.   For $n \in \mathbb{N}$, let $[n] = \{1, \dots, n\}$. Vectors are in bold, e.g., $\mathbf{a} = (a_1, \dots, a_\ell)$. Concatenation is $\|$. For a (possibly randomized) algorithm Alg, we write $y \leftarrow \mathsf{Alg}(x)$ for its output. The security parameter is λ, and $\mathrm{negl}(\lambda)$ denotes a negligible function.

2.1   Algebraic setting and basic primitives

Let $\mathbb{F}$ be a prime field of order q and $\mathbb{G}$ a cyclic group of order q with generator g, where discrete logarithms are hard. We use SHA-256 as $\mathsf{H} : \{0,1\}^* \to \{0,1\}^{256}$. The issuer uses a digital signature scheme $\mathsf{Sig}_I = (\mathsf{KeyGen}_I, \mathsf{Sign}_I, \mathsf{Verify}_I)$ (e.g., ECDSA/RSA), and the device uses Sig_D for nonce-bound session signatures; both are assumed EUF-CMA.

Pedersen vector commitments.   Fix generators $\mathbf{g} = (g_1, \dots, g_\ell)$ and h in $\mathbb{G}$. For $\mathbf{a} \in \mathbb{F}^\ell$ and $r \in \mathbb{F}$ define

$$\mathsf{Com}(\mathbf{a}; r) = \prod_{i=1}^{\ell} g_i^{a_i} \cdot h^r \in \mathbb{G}.$$

Com is perfectly hiding and binding under discrete-log hardness. We will need efficient *openings* and *equality proofs* for Pedersen commitments; we realize these via inner-product arguments (IPAs) in the Hyrax framework [WTs+18].

Hyrax commitments.   Hyrax commitments [WTs+18] allow us to commit to a multilinear extension $\widetilde{P}$ of a function $P : \{0,1\}^{\log n} \to \mathbb{F}$. Notably, we can express $\widetilde{P}(x_1, \ldots, x_{\log n})$ as a vector-matrix-vector product $\vec{v_L} P \vec{v_R}$, where:

- P is a square matrix of evaluations of P on the $\{0,1\}^{\log n}$ hypercube, represented as $\langle \vec{p_i} \rangle_{i \in \sqrt{n}}$ where $\vec{p_i} \in \mathbb{F}^{\sqrt{n}}$ are columns of P

- $\vec{v_L}(x_1, \ldots, x_{(\log n)/2}) = \langle \widetilde{\chi}_{(b_1, \ldots, b_{(\log n)/2})}(x_1, \ldots, x_{(\log n)/2}) \rangle_{(b_1, \ldots, b_{(\log n)/2}) \in \{0,1\}^{(\log n)/2}}$, and

- $\vec{v_R}(x_{(\log n)/2+1}, \ldots, x_{\log n}) = \langle \widetilde{\chi}_{(b_{(\log n)/2+1}, \ldots, b_{\log n})}(x_{(\log n)/2}, \ldots, x_{\log n}) \rangle_{(b_{(\log n)/2}, \ldots, b_{\log n}) \in \{0,1\}^{(\log n)/2}}$,

and $\widetilde{\chi}_{b_1, \ldots, b_m}(x_1, \ldots, x_m) = \prod_{i \in [m]} \widetilde{\chi}_{b_i}(x_i)$, where $\widetilde{\chi_{b_i}}(x_i)$ is the multilinear extension of the function $\chi_{b_i}(x_i) : \{0,1\} \to \mathbb{F}$ given by $\chi_{b_i}(x_i) = x_i b_i + (1 - x_i)(1 - b_i)$, which equals 1 if $x_i = b_i$ and 0 otherwise.

Then our Hyrax commitment scheme is given by:

- **Setup**$(1^\lambda, n) = (g_i)_{i \in [\sqrt{n}]}$ where $g_i \in \mathbb{G}$ are elements of the elliptic curve group $\mathbb{G}$ over which we compute our Pedersen commitments.

- **Commit**$_h(pp, P) \to (c, S)$, where $c = \{c_i\}_{i \in [\sqrt{n}]}$ and $S = \{S_i\}_{i \in [\sqrt{n}]}$, and $(c_i, S_i) \leftarrow Commit_p(pp, \langle \vec{p_i} \rangle)$ where $Commit_p$ is the Pedersen vector commitment scheme.

- **Eval**$_h(pp, c, r, v, n; P, S)$, where r is the opening point and v is the claimed evaluation of $P(r)$, is given by:

  – Prover and Verifier compute $\vec{v_L} = \vec{v_L}(r_1, \ldots, r_{(\log n)/2})$, $\vec{v_R} = \vec{v_R}(r_{(\log n)/2+1}, \ldots, r_{\log n})$, and $C = \sum_{i \in [\sqrt{n}]} \vec{v_L}_i c_i$ (where addition is over $\mathbb{G}$)

  – Prover and Verifier engage in an interactive IPA protocol to prove that $P(r)$ is a dot product of v_R and C

2.2   Zero-knowledge proving interface

We use a *transparent* general-purpose SNARK in the Spartan family [Set20], instantiated over $\mathbb{F}$ for R1CS instances. Algorithms are

$$\mathsf{pp} \leftarrow \mathsf{Setup}(1^\lambda), \quad \pi \leftarrow \mathsf{Prove}(\mathsf{pp}, x, w), \quad b \leftarrow \mathsf{Verify}(\mathsf{pp}, x, \pi) \in \{0, 1\},$$

where x is public input and w the witness. (We reference succinct CRS-based systems such as Groth16 [Gro16] later only for performance comparisons.)

R1CS instance.   An *R1CS instance* is a tuple $(\mathbb{F}, A, B, C, io, n, m)$, where io denotes the public input and output of the instance, $A, B, C \in \mathbb{F}^{n \times n}$, where $m \geq |io| + 1$ and there are at most m non-zero entries in each matrix. An instance is said to be *satisfiable* if there exists a witness $\vec{w} \in F^{n - |io| - 1}$ such that $(A \cdot Z) \circ (B \cdot Z) = (C \cdot Z)$, where $\vec{Z} = (io, 1, \vec{w})$, and $\circ$ is the Hadamard (entry-wise) product.

Throughout the paper, assume that we are dealing with sparse R1CS instances, where $m = O(n)$.

R1CS as a language   We let the language

$$R_{\mathrm{R1CS}} = \{\langle x = (\mathbb{F}, A, B, C, io, n, m) : x \text{ is satisfiable } \rangle\}.$$

The language R_{R1CS} is NP-complete.

A Spartan zero-knowledge proof is an argument of knowledge for the R1CS language. The Prover can prove not only that a given instance $x \in R_{\mathrm{R1CS}}$, but knowledge of the corresponding witness $\vec{w}$. At a high level, this application considers R1CS instances that represent the computational structure of ownership of a valid credential, along with any other desired properties about the credential. Thus a valid argument of knowledge implies knowledge of the underlying credential that has the claimed properties, which is necessary for real-world authentication.

2.3   High-Level Credential Presentation Flow

The following outlines the high-level protocol for a credential presentation.

1. The Prover receives a signed credential from an Issuer to be stored securely in their wallet, issued to the Wallet Secure Cryptographic Device (WSCD) public key.

2. At presentation time, the Verifier sends over challenge $\mathtt{nonce}_V$ for device-binding verification.

3. The Prover signs the challenge $\mathtt{nonce}$ with the public key p_U controlled by their WSCD and specified in their credential.

4. The Prover computes two separate but linked zero-knowledge proofs $\pi_{\mathtt{prepare}}, \pi_{\mathtt{show}}$ which together cover the following statements: SD-JWT parsing, verification of the SD-JWT Issuer signature, proper disclosures and/or arbitrary predicates on the disclosures, and device-binding. (i.e. checks the $\mathtt{nonce}$ signature against their public key); then sends $\pi_{\mathtt{prepare}}, \pi_{\mathtt{show}}$ to the Verifier.

5. The Verifier verifies $\pi_{\mathtt{prepare}}, \pi_{\mathtt{show}}$ independently, and also verify that they are linked; grants Prover access to some service based on their credential disclosures.

2.4   Security model

In our security model, we assume that the Prover is malicious, and that each Verifier is semi-honest, meaning that if the Prover presents a valid proof that they own a credential with some property, the Verifier will grant access to any services for which the property suffices.

For security on the Verifier's side, our soundness analysis considers the probability that a Prover without real ownership of a valid credential can generate a false proof of ownership.

For security on the Prover's side, we guarantee that our proofs are zero-knowledge, so that a semi-honest computationally bounded Verifier cannot get any additional private information about the Prover's credential given the proof, beyond what is publically revealed in the proof. In particular, we do not consider the case where the Verifier is malicious, where e.g. a false Verifier pretends to be an authorized Verifier. The problem of Verifier identity lies outside the scope of this paper.

Futhermore, we assume that Verifiers can collude, i.e. that Verifiers $V_1, \ldots, V_N$ that have receieved proofs $\{\pi_1\}, \ldots, \{\pi_N\}$ from a given Prover P can compute functions $f(\pi_1, \ldots, \pi_N)$. Therefore, we desire the **unlinkability property**: given $pi_1, \ldots, \pi_N$, the Verifiers should not be able to determine whether or not any two of these proofs came from the same Prover P. Note that this requires the Prover to re-randomize each presentation's proof; a static zero-knowledge proof of the same statement, while not revealing private credential information, will still look the same. It is possible that Verifier's can effectively de-anonymize a Prover by linking their anonymous activity across presentations and analyzing corresponding metadata, e.g. time of presentation.

3   Our zkID

At a high-level, we propose a generic zkSNARK wrapper over an EUDI digital credential, which will either be issued in SD-JWT format or the mDL data format specified in standard ISO/IEC 18013-5). The backend proving system we use will be a combination of Spartan and Hyrax commitments, with modifications to be zero-knowledge.

There are two (2) key ideas to highlight within our proposed architecture:

- **Pre-processing batches of re-randomized proofs** of issuer-signature and credential-parsing, to re-use across each new presentation. We call this the `prepare` relation.

- **Committing to the credential disclosures with Hyrax commitments** [WTas$^+$17], which allows us to re-use (or "link") witnesses across circuits "for free"

For the first item, we note that pre-processing proofs for the `prepare` relation is possible because the relation is independent of the presentation, including the choice of disclosures or predicate proofs. Further optimizations can likely be made to only parse the discosures of certain attributes within the credential if it is known that the Wallet User will rarely or never present certain disclosures to external verifiers.

The second items differs from the linking circuits approach that Google uses [Fas24]. Note that Google verifiably computes a hiding and binding MAC of the shared witnesses as a public output of the circuit, which the verifier checks consistency of in plain. Although this is only a few linear relations, it requires the prover to also commit to their portion of the key to prevent forging. We instead simply manually separate out the disclosures $m_1, \ldots, m_n$ into a designated column when committing to the witness, which is already needed to prove the circuit relation. Then, the verifier simply checks consistency of these commitments when verifying each circuit's proof.

We note that by using Hyrax commitments, our PoC is not post-quantum secure. In particular, post-quantum computers break the discrete-log assumption, which breaks the binding property of the polynomial commitment scheme. Thus, a malicious Prover could potentially make false proofs about their identity. In future work, we hope to incorporate modified Ajtai lattice-based commitments to ensure post-quantum security [HSS24].

Throughout the remainder of this section, we refer to the EUDI's Wallet User as the "Prover", the Relying Party as the "Verifier", and the EUDI Attestation Authority (EAA) as the "Issuer". Below, we briefly detail a high-level flow of the interaction between the Issuer, Prover, and Verifier.

3.1   Underlying ZK Circuit

In this section, we describe our high-level ZK circuit C underlying the knowledge the prover needs to present to the verifier. Throughout, we refer to the Issuer with variable I, Prover with variable P, and Verifier with variable V (e.g. in subscripts).

We will detail the ZK wrapper around the SD-JWT credential as an example, but the protocol is analogous for other credential formats. Throughout, we will refer to digests as "message hashes" or just "hashes", and disclosures as "messages".

We define a circuit C for proving ownership of an anonymous credential. We let our witness $w = S$ be the SD-JWT credential consisting of messages $\{m_i\}_{i=1}^N$, hash salts $\{s_i\}_{i=1}^N$, hashes $\{h_i\}_{i=1}^N$, and an Issuer signature $\sigma_I = \sigma(h_1, \ldots, h_N; SK_I)$. Without loss of generality, we assume that the Prover's public key PK_P is contained in message m_1 of the credential and indexable as $m_1[1]$. We let our instance $x = (PK_I, \{f_i\}_{i=1}^K, \{p_i\}_{i=1}^K)$ contain the Issuer's public key PK_I, functions f_i over the messages, output predicates p_i, and finally the nonce signature σ_{nonce} for proving device-binding. The f_i can be arbitrary

statements we wish to prove about the messages. For example, one could define a function $f_i(m_1, \ldots, m_N) = m_1$ would output a predicate that is just the disclosure of message m_1.

Underlying ZK Circuit C for Verifiable Credential

We define circuit $C(x = (PK_I, \{f_i\}_{i=1}^K, \{p_i\}_{i=1}^K), w_C = (S))$ as follows:

1. Assert $\mathrm{parse}_{\mathrm{SD\text{-}JWT}}(S) = (\{m_i\}, \{s_i\}, \{h_i\}_i, \sigma_I)$ parsing of the SD-JWT into messages $\{m_i\}_{i=1}^N$, message salts $\{s_i\}_{i=1}^N$, hashes $\{h_i\}_{i=1}^N$ and Issuer signature σ_I.

2. Assert $h_i = \mathrm{SHA256}(m_i, s_i) \quad \forall i \in [n]$, i.e. that messages hashes correspond to messages and salts

3. Assert $p_i = f_i(m_1, \ldots, m_n) \quad \forall i \in [n]$, i.e. correct evaluation of the predicates

4. Assert $\mathrm{ECDSA.verify}(\sigma_I, PK_I) = 1$, i.e. the credential signature verifies under the Issuer public key

5. Assert $\mathrm{ECDSA.verify}(\sigma_{\mathrm{nonce}}, m_1[1]) = 1$, i.e. that the live nonce signature corresponds to the public key the credential was issued to

---

3.2 Pre-processing and linking proofs

The main speedups from our proving system will come from splitting our high-level circuit C above into two (2) circuits for different relations regarding the digital credential, namely a `prepare` and a `show` relation, analogous to Microsoft's Crescent Credentials [PPZ24]. This is advantageous because proofs of the `prepare` relation can be computed a-priori for any credential, as they do not depend on the claim being proved at presentation time. Pre-computing these proofs will save significant time per presentation, and reduce the performance bottleneck to that of proving the `show` relation.

One issue that arises is the need to ensure consistency of witnesses across these separate circuits, or what we call "linking proofs". At a high level, as opposed to Google's MAC approach [Fas24], the prover sends Hyrax commitments to the parts of the witness re-used across circuits, which ends up being just the raw messages $\{m_i\}_i$. The verifier can then check consistency of these witnesses across the circuits C_i by comparing the Hyrax commitments they receive as part of the proof. This approach gives us linking "for free", as the Prover already needs to compute these Hyrax commitments as part of the proof.

We highlight that we are no longer splitting up circuits by their field operations (e.g. SHA256 attestations over a binary extension field and an ECDSA verifications over a prime field), but rather by pre-processing and per-presentation relations. In particular, the circuit for `prepare` will necessarily involve wrong-field arithmetic by including both the SHA256 hashes and the Issuer ECDSA signature verification. However, because of the ability to pre-compute proofs of the `prepare` relation, the more important thing becomes to choose curves that allow for i) efficient show relations and ii) linking the prepare and show relation. Since the verifier can only check equality of Hyrax Pedersen commitments defined over the same curve, we must use the same curve for proving both the `prepare` and `show` relations. Thus we choose a curve with a scalar field equivalent to the base field of the nonce signature σ_{nonce} for efficient signature verification. Because most Hardware Security Modules (HSMs) sign over the P256 curve, we choose the Tom256 (T256) curve for our backend, which has scalar field equivalent to the base field of P256.

We now detail each of the two (2) relations/circuits below.

3.2.1 The `prepare` relation:

The `prepare` relation checks the validity of issuer signature, parses the SD-JWT, and verifies all the message hashes, none of which depend on the specific presentation. Thus, the prover will periodically pre-compute and store a batch of re-randomized proofs of the prepare relation. These proofs will utilize Hyrax Pedersen vector commitments as introduced above in order to link the proofs of `prepare` relation to the `show`.

Circuit C_1 for the `prepare` relation

We define circuit $C(x = (PK_I), w_i = S, w = (\{m_i\}_{i=1}^N))$ as follows:

1. Assert $\mathrm{parse}_{\mathrm{SD\text{-}JWT}}(S) = (\{m_i\}, \{s_i\}, \{h_i\}_i, \sigma_I)$ parsing of the SD-JWT into messages $\{m_i\}_{i=1}^N$, message salts $\{s_i\}_{i=1}^N$, hashes $\{h_i\}_{i=1}^N$ and Issuer signature σ_I.

2. Assert $h_i = \mathrm{SHA256}(m_i, s_i) \quad \forall i \in [n]$, i.e. that messages hashes correspond to messages and salts

3. Assert $\mathrm{ECDSA.verify}(\sigma_I, PK_I) = 1$, i.e. the credential signature verifies under the Issuer public key

---

The backend proving system we will use for verifiably computing circuits is Spartan, coupled with a Hyrax-style Pedersen commitment scheme. We can express the circuit computation as some R1CS relation

$$(A \cdot Z) \circ (B \cdot Z) = (C \cdot Z),$$

where $\vec{Z} = (io, 1, \vec{w})$ and io are the public input/outputs. Spartan proves knowledge of a vector Z of length $n := |Z|$ that satisfies the R1CS instance.

To produce zkSNARK proofs for this circuit C_1, the prover will proceed in two phases:

1. `prepareCommit`: Separates out a column containing only message hashes $\{m_i\}_{i \in [N]}$ in Z and computes an initial Hyrax commitment $c^{(1)} = \{c_i^{(1)}\}_{i \in [\sqrt{n}]}$, which includes a Pedersen commitment to the messages column $c_1^{(1)} = com(m_1, \dots, m_N; r_1^{(1)}) = g_1^{m_1} \dots g_N^{m_N} g_{N+1}^{r_1^{(1)}}$ with initial randomness $r_1^{(1)}$.

2. `prepareBatch`:

   (a) Re-randomizes this initial Hyrax commitment to get a batch of commitments $c^{(j)} = \{c_i^{(j)}\}_{i \in [\sqrt{n}]}$, each of which contains a Pedersen commitment to the messages $c_1^{(j)} = com_1^{(1)} \cdot g_{N+1}^{r_1^{(j)} - r_1^{(1)}}$ for all $j \in [m]$, where our batch size m depends on the frequency of proof generation and demand for the credential

   (b) Continues the Spartan sumcheck IOP on each $c^{(j)}$ to produce a batch of proofs $\{\pi_{\mathrm{prepare}}^{(j)}\}$ for $j \in [m]$ of the `prepare` relation.

The prover will run `prepareBatch` periodically to both generate re-randomized commitments $c^{(j)}$ and store the randomness for the message column commitment $r_1^{(j)}$ for linking purposes, as well as generate and store batches of issuer-signature proofs $\pi_{\mathrm{prepare}}^{(j)}$ that can be used for each presentation.

3.2.2 The `show` relation:

At a high-level, our show relation will i) verifiably compute any functions f_i over the SD-JWT messages (such as disclosures, range checks, etc.), and ii) check that the credential

belongs to the prover's device (also known as proof of "device-binding"). As part of device-binding, the prover will sign a verifier `nonce` outside of the circuit, as outlined in flow 2.3. Let us denote this signature by $\sigma_P = \sigma(\text{nonce}; SK_P)$

Again, we will use T256 curve for our backend proving system so that the holder in-circuit signature verification can proceed naturally in the right field.

Circuit C_2 for the `show` relation

We define circuit $C_2(x = (\{f_i\}_{i=1}^K, \{p_i\}_{i=1}^K), w = \{m_i\}_{i=1}^N)$ as follows:

1. Assert $p_i = f_i(m_1, \ldots, m_n) \quad \forall i \in [n]$, i.e. correct evaluation of the predicates

2. Assert $\text{ECDSA.verify}(\sigma_{\text{nonce}}, m_1[1]) = 1$, i.e. that the live nonce signature corresponds to the public key the credential was issued to

---

As part of computing proof $\pi_{\text{show}}^{(j)}$ for presentation $j \in [m]$, the Prover will once again separate out the messages into a separate column to compute a Hyrax commitment over the Tom256 curve. In particular, the Prover uses *the same* randomness $r_1^{(j)}$ used during the `prepareBatch` process to compute the Pedersen commitment to the messages column. The verifier will then check that the Pedersen commitment to the messages column for $\pi_{\text{show}}^{(j)}$ equals that of proof $\pi_{\text{prepare}}^{(j)}$ for circuit C_1.

3.3 Adding ZK to Spartan

Our construction uses Circom in the frontend to compile our computation into an R1CS (instance, witness) pair $(x = (\mathbb{F}, A, B, C, io, n, m), \vec{w})$, which we then feed into the Spartan IOP coupled with Hyrax-style Pedersen polynomial commitments.

Recall that our R1CS constraint looks like the following:

$$(A \cdot \vec{Z}) \circ (B \cdot \vec{Z}) - (C \cdot \vec{Z}) = 0$$

where our square matrices A, B, C have size n and $\vec{Z} = (\vec{w}, 1, io)$.

Recall that Spartan converts an R1CS constraint into the following zero-check:

$$\sum_{x \in \{0,1\}^{\log n}} \widetilde{eq}(x, \tau) \left[\left(\sum_{y \in \{0,1\}^{\log n}} \widetilde{A}(x, y) \widetilde{Z}(y) \right) \left(\sum_{y \in \{0,1\}^{\log n}} \widetilde{B}(x, y) \widetilde{Z}(y) \right) \right.$$
$$\left. - \left(\sum_{y \in \{0,1\}^{\log n}} \widetilde{C}(x, y) \widetilde{Z}(y) \right) \right] = 0$$

for some random challenge $\tau \in \mathbb{F}$

There are two components in Spartan that we need to modify to be ZK. The first is making the sumchecks ZK. The second is to ensure that the opening $\widetilde{Z}$ using the commitment to $\vec{Z}$ does not leak information about our witness $\vec{w}$.

3.3.1 Adding ZK to sumcheck

The Spartan protocol consists of several sumchecks in parallel and operates over some field $\mathbb{F}$. There are various existing techniques to make sumcheck ZK. We employ one using similar methods as in Zhang et. al. [ZXZS19], which adds uniformly random pads to the sumcheck transcript.

In particular, suppose at each round i of the sumcheck protocol, the prover sends over $s_i(X) := \sum F(r_1, \ldots, r_{i-1}, X, x_{i+1}, \ldots, x_m)$ where r_i is the Verifier challenge sent for round i. Then instead of having the verifier check the sumcheck, the prover will prove in

ZK that the unpadded transcript satisfies the verifier's (linear) checks. To do this, the prover will need to commit to the uniformly random pads ahead of time. Then, as long as the Fiat-shamir challenges is generated from the transcript including these random pad commitments, the prover cannot simply lie about the pads to satisfy the sumcheck relation.

Adding ZK to sumcheck

1. Prover commits to pads $R_i(X) \xleftarrow{\$} \mathbb{F}_1[x]$ for all $i \in [\log n]$. These are linear polynomials, and can thus be represented by its two coefficients $R_i[0]$ and $R_i[1]$.

2. Instead of sending partial sums

$$s_i(X) := \sum_{(x_{i+1},\ldots,x_n) \in \{0,1\}^{n-i}} F(r_1,\ldots,r_{i-1},X,x_{i+1},\ldots,x_n)$$

   for each round of sumcheck, the Prover sends polys $s_i'(X) = s_i(X) + R_i(X)$, essentially a one-time-padded transcript.

3. It suffices to show the following linear relation in zero-knowledge,

$$\begin{bmatrix} A \mid B \mid C \end{bmatrix} \begin{bmatrix} \vec{S} \\ \hline \vec{R} \\ \hline C \\ F(r) \end{bmatrix} = \vec{0}$$

   where
$$\vec{S} = \begin{bmatrix} s_1'[0], s_1'[1], \ldots, s_n'[0], s_n'[1] \end{bmatrix}^\mathsf{T},$$
   is the column vector of sumcheck transcripts such that $s_i' = s_i'(X) = s_i'[0]X + s_i'[1]$, and
$$\vec{R} = \begin{bmatrix} R_1[0], R_1[1], \ldots, R_n[0], R_n[1] \end{bmatrix}^\mathsf{T},$$
   is the column vector of random pads, and where matrices A, B, C are given by

$$A = \begin{bmatrix}
1 & 2 & 0 & 0 & 0 & 0 & 0 & \cdots & 0 & 0 & 0 & 0 \\
-r_1 & -1 & 1 & 2 & 0 & 0 & 0 & \cdots & 0 & 0 & 0 & 0 \\
0 & 0 & -r_2 & -1 & 1 & 2 & 0 & \cdots & 0 & 0 & 0 & 0 \\
\vdots & \vdots & \vdots & \vdots & \vdots & \vdots & \vdots & \ddots & \vdots & \vdots & \vdots & \vdots \\
0 & 0 & 0 & 0 & 0 & 0 & 0 & \cdots & -r_{n-1} & -1 & 1 & 2 \\
0 & 0 & 0 & 0 & 0 & 0 & 0 & \cdots & 0 & 0 & r_n & 1
\end{bmatrix}$$

$$B = \begin{bmatrix}
-1 & -2 & 0 & 0 & 0 & 0 & 0 & \cdots & 0 & 0 & 0 & 0 \\
r_1 & 1 & -1 & -2 & 0 & 0 & 0 & \cdots & 0 & 0 & 0 & 0 \\
0 & 0 & r_2 & 1 & -1 & -2 & 0 & \cdots & 0 & 0 & 0 & 0 \\
\vdots & \vdots & \vdots & \vdots & \vdots & \vdots & \vdots & \ddots & \vdots & \vdots & \vdots & \vdots \\
0 & 0 & 0 & 0 & 0 & 0 & 0 & \cdots & r_{n-1} & 1 & -1 & -2 \\
0 & 0 & 0 & 0 & 0 & 0 & 0 & \cdots & 0 & 0 & -r_n & -1
\end{bmatrix}$$

$$C = \begin{bmatrix} -1 & 0 \\ 0 & 0 \\ 0 & 0 \\ \vdots & \vdots \\ 0 & 0 \\ 0 & -1 \end{bmatrix}$$

4. The Prover computes a public random challenge α (e.g. hashing the transcript) and compresses the relation into a dot product

$$[\vec{u}] \begin{bmatrix} \vec{S} \\ \hline \vec{R} \\ \hline C \\ F(r) \end{bmatrix} = \vec{0} \tag{1}$$

where $\vec{u} = [1, \alpha, \alpha^2, \ldots, \alpha^n] \begin{bmatrix} A \mid B \mid C \end{bmatrix}$ is a random linear combination of the rows.

5. The Prover and Verifier engage in a proof-of-dot product protocol to prove the relation above, such as an Inner Product Argument used in Bulletproofs [BBB$^+$17]

3.3.2 Adding ZK to the opening of $\widetilde{Z}$

In order to add ZK to the opening of, we simply "append" Z with uniformly random pads. Specifically, we assign random evaluations $Z(x) \xleftarrow{\$} \mathbb{F}$ on any remaining point in the hypercube $x \in \{0,1\}^{\log n}$. Then, we see that $\widetilde{Z'}(x_1, \ldots, x_{\log n}) = \widetilde{Z}(x_1, \ldots, x_{\log n}) + eq(r, x) Z(x)$. Note that now $\widetilde{Z'}(x_1, \ldots, x_{\log n})$ is distributed uniformly at random. If $n = |Z|$ is not already a power of 2, then we can simply fill at least one of the remaining evaluations on $\{0,1\}^{\log n}$ with a single random pad. If $n = 2^m$, we can add another dimension to the hypercube of evaluations of Z.

3.4 Cost analysis

The following section computes the Prover and Verifier costs of Spartan instantiated with Hyrax Pedersen commitments on R1CS instances $(x = (\mathbb{F}, A, B, C, io, n, m), \vec{w})$, where io denotes the vector of public inputs/outputs, $n = |\vec{w}| + io + 1$ is the dimension of our matrices, and m is the number of nonzero entries in our matrices A, B, C. We let $Z = (\vec{w}, io, 1)$. It is often reasonable to assume that our R1CS matrices are sparse, i.e. $m = O(n)$. However, we present the costs below independent of this assumption.

- **Prover time**: (1) $O(m)$ to generate sumcheck transcript, (2) $O(m)$ to evaluate MLEs of A, B, C, (3) $O(n)$ to commit to the MLE of Z (computing $\sqrt{n}$ MSMs of size $\sqrt{n}$) and opening the MLE of Z, for a total cost of $O(m)$.

- **Proof length**: (1) $O(\log n) \cdot |\mathbb{F}|$ length of the sumcheck transcript, (2) $O(\sqrt{n}) \cdot |\mathbb{G}|$ length commitment to MLE of Z, (3) $O(\log n) \cdot |\mathbb{G}|$ length of argument opening MLE of Z, for a total length of $O(\sqrt{n})$ group or field elements.

- **Verifier time**: (1) $O(\log n)$ to verify sumcheck transcript, (2) $O(m)$ to evaluate the MLEs of A, B, C (with sparse commitment scheme and memory checking), (3) $O(\sqrt{n})$ to open the MLE of Z, for a total of $O(m + \sqrt{n})$.

With the ZK modifications to Spartan, we can see the asymptotic costs remain the same, as follows:

- **Prover time**: additionally computes $O(\log n)$ constant-size commitments to the sumcheck transcript pads r_i, and $O(\log n)$ engages in new sumcheck relation IPA (or some other ZK dot product argument) for vector of length $O(\log n)$, which still gives $O(m)$ prover work.

- **Proof length**: sumcheck and openings are the same length but just padded, but added on $O(\log n)$ size $|\mathbb{G}|$ commitments, and a length $O(\log \log n)$ sumcheck relation IPA proof, which still gives a proof length of $O(\sqrt{n})$ group or field elements.

- **Verifier time**: no longer needs to do $O(\log n)$ (sumcheck), but still needs $O(m)$ (evaluating MLEs of A, B, C) + $O(\sqrt{n})$ (opening MLE of Z) + $O(\log n)$ for sumcheck relation IPA verification, which still gives $O(m + \sqrt{n})$ runtime.

3.5   Security analysis

The correctness follows immediately from the correctness of the Spartan SNARK and the fact that the Prover uses the same randomness for the Hyrax commitments across the `show` and `prepare` circuits for each presentation i.

The soundness of our protocol follows from the soundness of Spartan. In particular, we can extract the full witness credential from the `prepare` relation.

Intuitively, zero-knowledge follows from the hiding property of the commitment scheme as well as the zero knowledge property of the Spartan zkSNARK proving system; For proof i, simulator can randomly sample the linked commitment com_i both distributions to reuse across both proofs, both in the commitment itself and also in the IPA used to open the Hyrax commitment to $Z(r_1, \ldots, r_{\log n})$. We can show that this commitment is independent of the rest of the view of the Verifier, which consists of the following:

- Sumcheck polynomials $\{s_i'(X)\}_{i \in [\log n]}$ for each of the sumchecks in Spartan

- $\{r_i\}$ Fiat-Shamir challenges during the sumcheck

- Transcript from the IPA on the sumcheck relation in ZK

- $\{com(z_i)\}_{i \in [\sqrt{n}]}$ Hyrax commitment to Z, which involves a Pedersen commitment to each of the columns of a $\sqrt{n} \times \sqrt{n}$ matrix representation of $\vec{Z}$

- Transcript from the IPA for opening $Z(r_1, \ldots, r_{\log n})$

- The claimed value of $Z(r_1, \ldots, r_{\log n})$

Since we appended random pads to $\vec{Z}$ in our ZK modifiction in Section 3.3.2, the distribution of $Z(r_1, \ldots, r_{\log n})$ is random and independent of Z, and therefore independent of $\{m_i\}_i$. Furthermore, $s_i'(X)$ have totally random pads on them and their distribution is independent of Z, and therefore independent of $\{m_i\}_i$. Assuming the hiding property of the Pedersen commitment schemes for messages sent during an IPA, we can also use the simulators for the IPAs without changing their joint distribution with the rest of the transcript.

Then, we can simply run the piece-wise simulators for each zkSNARK proof for circuits C_1 and C_2 to simulate the remainder of the view.

4 Experiments

5 Application to EUDI

6 Security

In our security model, we assume that the Prover is malicious, and that each Verifier is semi-honest, meaning that if the Prover presents a valid proof that they own a credential with some property, the Verifier will grant access to any services for which the property suffices.

Verifier's side For security on the Verifier's side, our soundness analysis considers the probability that a malicious Prover without real ownership of a valid credential can generate a false proof of ownership.

Prover's side For security on the Prover's side, we guarantee that our proofs are zero-knowledge, so that a semi-honest and computationally-bounded Verifier cannot get any additional information about the Prover's credential beyond what is publically revealed in the proof. In particular, we do not consider the case where the Verifier is malicious during presentation, e.g. where a false Verifier pretends to be an authorized Verifier. The problem of Verifier identity lies outside the scope of this paper.

Furthermore, we assume that Verifiers can collude with each other, i.e. that Verifiers $V_1, \ldots, V_N$ that have received proofs $\{\pi_1\}, \ldots, \{\pi_N\}$ from a given Prover P can compute functions $f(\pi_1, \ldots, \pi_N)$. Therefore, we desire the **unlinkability property**: given $pi_1, \ldots, \pi_N$, the Verifiers should not be able to determine whether or not any two of these proofs came from the same Prover P. Note that this requires the Prover to re-randomize each presentation's proof; a static zero-knowledge proof of the same statement, while not revealing private credential information, would still look the same across presentations. In that case, it may be possible for the Verifier to de-anonymize a Prover by linking their "anonymous" activity across presentations and analyzing metadata, e.g. time of presentation. Fortunately, our scheme is unlinkable due to the re-randomization of proofs between each presentation. By the zero-knowledge property for each presentation, we can simulate the distribution of proofs without knowledge of the witness. To simulate an entire set of proofs received by distinct colluding Verifiers, we can independently simulate each proof.

Finally, as our scheme is currently presented in Section 3, we assume that Verifiers will not collude with Issuers even though they can see the Issuer public key. To bypass this assumption and prevent Issuer tracking in the case of malicious Verifiers that collude with Issuers, we propose here the maintenance of a trusted Merkle tree on trusted Issuer public keys. Then our Prover's circuit would prove knowledge of a valid Issuer signature from some key in the Merkle tree, and the public input/output would just be the Merkle root rather than any specific Issuer public key.

Both Prover and Verifier security When modelling the Issuer, we assume that the Issuer is trusted during issuance by both the Prover and Verifier, i.e. will not Issue false credentials or sell personal information that is necessarily to obtain about individuals to issue a credential.

6.1 Other considerations

Our scheme currently does not require any interaction from the Issuer for credential presentation beyond initial issuance.

However, our scheme does require the use of internet access (without, there will be risks with authorizing someone before their credential can be checked against the current state). In the case (as presented) where Issuer public key is a public input/output, we assume there is an online registry of trusted Issuer keys that the Verifier can check the proof against. This requires live internet access in the same way that credit card transactions do, in order to check the most current registry of public keys. Even in the case of a Merkle inclusion proof, where the Issuer key is also private, the Verifier would need to check that the public Merkle root matches the trusted root stored online. It is possible to store encrypted transactions/credential presentations to be checked later once internet access is restored, in the same way as offline credit card transactions do. However, there are necessary risks with this approach; it would be up to the specific vendor and/or service prover what levels of risk can be tolerated from delayed credential authentication. For example, some service provider (Verifier) may be fine only periodically downloading the current registry of trusted Issuer keys (and/or Merkle roots) and simply checking against their last downloaded version before granting access.

Finally, as mentioned previously, our scheme is not quantum resistant due to the use of Hyrax commitments. Again, we believe this is easily fixable with the introduction of modified Ajtai lattice-based commitments, which are post-quantum secure.

7 Conclusion

References

[AHIV17]   Scott Ames, Carmit Hazay, Yuval Ishai, and Muthuramakrishnan Venkitasubramaniam. Ligero: Lightweight sublinear arguments without a trusted setup. In Bhavani M. Thuraisingham, David Evans, Tal Malkin, and Dongyan Xu, editors, *ACM CCS 2017: 24th Conference on Computer and Communications Security*, pages 2087–2104, Dallas, TX, USA, October 31 – November 2, 2017. ACM Press.

[BBB⁺17]   Benedikt Bünz, Jonathan Bootle, Dan Boneh, Andrew Poelstra, Pieter Wuille, and Greg Maxwell. Bulletproofs: Short proofs for confidential transactions and more. Cryptology ePrint Archive, Paper 2017/1066, 2017.

[Fas24]   Matteo Frigo and abhi shelat. Anonymous credentials from ECDSA. Cryptology ePrint Archive, Paper 2024/2010, 2024.

[FYC25]   Daniel Fett, Kristina Yasuda, and Brian Campbell. Selective disclosure for JWTs (SD-JWT). Technical Report draft-ietf-oauth-selective-disclosure-jwt-22, IETF OAuth WG, May 2025. Internet-Draft.

[Gro16]   Jens Groth. On the size of pairing-based non-interactive arguments. In *EUROCRYPT 2016*, volume 9665 of *LNCS*, pages 305–326. Springer, 2016.

[GRS⁺]   Paul A. Grassi, Justin P. Richer, Sarah K. Squire, James L. Fenton, and Ellen M. Nadeau. Digital identity guidelines: Federation and assertions (nist special publication 800-63c). NIST Special Publication 800-63C SP 800-63C, National Institute of Standards and Technology (NIST). Includes updates as of 2020-03-02; Privacy Authors: Naomi B. Lefkovitz, Jamie M. Danker; Usability Authors: Yee-Yin Choong, Kristen K. Greene, Mary F. Theofanos.

[HSS24]   Intak Hwang, Jinyeong Seo, and Yongsoo Song. Concretely efficient lattice-based polynomial commitment from standard assumptions. Cryptology ePrint Archive, Paper 2024/306, 2024.

[PPZ24]   Christian Paquin, Guru-Vamsi Policharla, and Greg Zaverucha. Crescent: Stronger privacy for existing credentials. Cryptology ePrint Archive, Paper 2024/2013, 2024.

[Set20]   Srinath Setty. Spartan: Efficient and general-purpose zkSNARKs without trusted setup. In Daniele Micciancio and Thomas Ristenpart, editors, *Advances in Cryptology – CRYPTO 2020, Part III*, volume 12172 of *Lecture Notes in Computer Science*, pages 704–737, Santa Barbara, CA, USA, August 17–21, 2020.

[Wor25]   World Wide Web Consortium (W3C). Verifiable credentials data model v2.0. W3C Recommendation REC-vc-data-model-2.0, World Wide Web Consortium (W3C), W3C, May 2025. Former editors: Grant Noble, Dave Longley, Daniel C. Burnett, Brent Zundel, Kyle Den Hartog.

[WTas+17] Riad S. Wahby, Ioanna Tzialla, abhi shelat, Justin Thaler, and Michael Walfish. Doubly-efficient zkSNARKs without trusted setup. Cryptology ePrint Archive, Paper 2017/1132, 2017.

[WTs+18]   Riad S. Wahby, Ioanna Tzialla, abhi shelat, Justin Thaler, and Michael Walfish. Doubly-efficient zkSNARKs without trusted setup. In *2018 IEEE Symposium on Security and Privacy*, pages 926–943, San Francisco, CA, USA, May 21–23, 2018. IEEE Computer Society Press.

[ZXZS19]   Jiaheng Zhang, Tiancheng Xie, Yupeng Zhang, and Dawn Song. Transparent polynomial delegation and its applications to zero knowledge proof. Cryptology ePrint Archive, Paper 2019/1482, 2019.

8   Appendix: EUDI Annex 2 Requirements

This secion is devoted to a review of the EUDI ARF's Annex 2, which covers high-level requirements for the EUDI Wallet. The full Annex can be found here.

In this Appendix, we address all of the Topics presented in Annex 2 of the EUDI Architecture Reference Framework. The first part of this section is addressed to a general audience; it translates the vocabulary used in the EUDI ARF to cryptographic terms that those reading this paper might be familiar with.

We then present some additional directions for our proof-of-concept that introduces potential solutions to some of the Topics (requirements) in the EUDI ARF. Finally, we lay out an exhaustive table addressing each of the Topics of the ARF, indicating: (i) whether or not our paper addresses and satisfies the specifications (see "Within Scope" column), (ii) any potential privacy concerns that may arise from satisfying the specification of the Topic (iii) a brief high-level proposals for any privacy concerns that do arise.

Topic	Summary	Comments	Within Scope	Paper Notes	Privacy Issues	Privacy Notes
1	Device binding and remote flows, according to [OpenID4VP] standard	Some of the points (e.g. verifying relying party identity, UX flows for which credential to present) are not strictly covered within our flow	Yes	Device binding: we propose that WSCA signature of nonce should be checked against pk of credential PRIVATELY in-circuit (otherwise reveals pk → linkability). Relying party verifying PID/QEAA signatures from trust list: can check issuer pk from list while hiding issuer signature check PRIVATELY in-circuit	Yes	We would propose merkle inclusion proof to hide the specific issuer pk (in case that issuer has only issued a few credentials), but would rely on external trusted maintenance/agreement of the merkle tree. This may not strictly align with "validate signature using trust list".
2	Wallet must support mDLs		Yes	Must support mDLs (specified in [ISO/IEC 18013-5]). Crescent already supports mDLs CBOR parsing	Maybe	
3	PID rulebook	Issuer responsibility	No		No	
4	mDL rulebook	Issuer responsibility	No		No	
5	EUDI wallet design	No HLRs	N/A		N/A	
6	Relying party authentication	User checks relying party ID and certificates, e.g. signature checks	No			
7	Only issuers can revoke, using an "attestation status/revocation list mechanism", that relying parties also use	Either short term credentials, have an attestation status (e.g. suspension) list mechanism, or attestation revocation list mechanism	Maybe	Prover needs to provide some kind of proof of non-revocation. We propose providing merkle inclusion proof of a public online list	Yes	Without attestation lists being public, would either i) need to phone home to issuer to see the status of credential/obtain a proof → issuer surveillance, or ii) provide an ID the relying party can check against a public list → linkability
8	N/A	No HLRs	N/A		N/A	
9	Wallet provider provides WTE certificate to wallet instances (testifying security of WSCA/D), wallet gives new PK and proof-of-association with WTE key to issuers to receive credentials to	Wallet provider responsibility, WTE is never shown to relying parties (see Topic 18 for proof-of-association sent to relying parties)	No		Yes	Potential issuer collusion allows for reconstruction of a superset of IDs if WTE is provided in plain. Proof-of-association with keys should also be private/in-circuit if possible
10	Wallets must support proximity and remote [OpenID4VP] flows. Wallets must support mDLs and SD-JWTs. UX flow around user accepting newly issued PID/attestation		Yes	Again must support mDLs (specified in [ISO/IEC 18013-5]). Crescent already supports mDLs CBOR parsing		

Topic	Summary	Comments	Paper Scope?	Paper Notes	Privacy Concerns?	Privacy Notes
11	Pseudonyms issued by a pseudonym provider. Allows relying party to recognize users across presentations		Yes	To eliminate need for an external pseudonym provider + allow for multiple pseudonyms controlled by the user: we propose computing/outputting a deterministic nullifier hash H(public_key, random_salt) (where the wallet stores the random_salt) in-circuit, to use as the pseudonym	Yes	If pseudonym provider is issuer, this is really bad (can track full identity whenever issuer-assigned pseudonym used). If pseudonym provider is an external party, perhaps need some kind of id disclosure to get a pseudonym → similar to issuer tracking
12	Should standardize attribute identifiers/syntaxes across namespaces and attestation types for max interoperability	Issuer responsibility / for verifier request convenience. But also somewhat for user legibility/transparency	No		No	
13		No HLRs	N/A		N/A	
14		No HLRs	N/A		N/A	
15		No HLRs	N/A		N/A	
16	Wallet should allow user to create (qualified electronic) signatures over documents		No		No	
17	Ensuring users with multiple accounts/credential signins are actually the same person — requesting linking IDs		Yes	Credentials should be re-randomized always	Yes	"Request the identified EUDI Wallet User to identify with another eID means which is accepted by the Relying Party so to link the data received from the EUDI Wallet with the account to which the User proved to have access to" → further supports need for efficient re-randomization
18	User presenting info/proofs across multiple credentials. Shall request proof-of-association of PKs each credential is issued to from the WCSA/WCSD	Here is a potential Schnorr-style ZKP for association between PKs. Main idea: proof of knowledge of dlog relationship btw the two	Yes	Proof-of-association should not have PKs in plaintext to the verifier (otherwise provides linkability). Need wallet user to do process proof received directly from WSCA. We propose an in-circuit verification of the ZKP received from WSCA (as a recursive proof) + matching of PKs to the ones in the credential	Yes	««

Topic	Summary	Comments	Paper Scope?	Paper Notes	Privacy Concerns?	Privacy Notes
19	Must have overview of all transactions executed through the Wallet Instance that cannot be deleted...		No		Yes	Need to ensure this is hidden behind some kind of biometric authentication within the wallet, so that cannot be accessed if stolen. Access to this data allows linkage everywhere. Definitely should not be external
20	ID verification should be strict to prevent fraud	No HLRs	N/A		N/A	
21		No HLRs	N/A		N/A	
22		No HLRs	N/A		N/A	
23		No new HLRs	N/A		N/A	
24	UX flows for proximity (e.g. mDL) - including user approval for disclosure	Just UX stuff	No		No	
25	Similar to topic 12, standardizing vocab for attestation attributes	Issuer responsibility	No		No	
26	For these standards, anyone can contribute, just be reasonable	Issuer/standards responsibility	No		No	
27	Relying party should get certificates to be verified by the wallet user before disclosure	Relying party + trusted anchor responsibility	No		No	
28	Wallet for legal person (e.g. corps, governments, and NGOs) should be diff from natural person (human)	Issuer responsibility	No		No	
29	Should allow for issuance of eIDs that let someone represent someone else. Not speced out yet, ad-hoc	Unclear	Maybe	We propose including both entities' PKs in the credential where transaction log would show up in represented person's asw. Maybe phone home to person being represented is OK?	Maybe	««
30	Wallet user can also be verifier	Nothing new	No		No	
31	Public, no-auth trusted list of certificate issuers (for issuers, wallet providers, and relying parties)	Nothing concerning! Public list is good	No		No	
32	PID interop	No HLRs	N/A		N/A	
33	Backup devices (e.g. HSM backups)		No		Yes	How do HSM backups work? How to ensure they are secure so that credentials aren't stolen

Topic	Summary	Comments	Paper Scope?	Paper Notes	Privacy Concerns?	Privacy Notes
34	Transfer to diff wallet	No HLRs	N/A		N/A	
35	Protocol for PID issuance	Yes	Yes	See privacy concerns around wallet user providing WTE for PID/attestation issuer to verify before credential issuance	Yes	In the cases that the issuer is just another verifier or EUDI wallet instance (user) — want WTE verification to be private (especially for credential issuances that don't require sensitive info/verification) for unlinkability
36	Risk analysis	No HLRs	N/A		N/A	
37	Remote signing	No HLRs	N/A		N/A	
38	Need to ensure use of wallet instance attestations (WIA) to relying parties -/→ tracking. Wallets must revoke/suspend when PID issuer asks		Yes		Yes	If WIAs are a separate list and just a static watermark, then can do merkle inclusion proof for unlinkability. We recommend against publishing list of public keys as a WIA mechanism — since this causes device-binding to function as doxxing (can check signatures against each of the potential public keys)
39	Wallet to wallet technical topic	No HLRs	N/A		N/A	
40	Reserved	No HLRs	N/A		N/A	
41	Min reqs on PuB-EAAs rulebooks	No HLRs	N/A		N/A	
42	Qualified trust service providers (qtsps)(issuers of QEAA attestation types) need access to authentic sources (e.g. databases of legal citizens)		No		No	
43	Policy for which items to be disclosed in certain presentation situations		No	Needs to be built into the frontend logic before circuit proofs on the user side.	Maybe	Only if some metadata about Wallet (that would provide linkability) can be tracked when requesting info from relying party...
44	QEAA eval reqs	Deleted	N/A		N/A	
45	QEAA rulebook	No new HLRs	N/A		N/A	
46	Protocol of cred presentation w verifiers	No HLRs	N/A		N/A	
47	Including non-QEAAs	No HLRs	N/A		N/A	
48	Users being able to delete presented data to verifiers		No		Yes	Good!
49	Protocols for 48	Deleted	N/A		N/A	

Topic	Summary	Comments	Paper Scope?	Paper Notes	Privacy Con-cerns?	Privacy Notes
50	Protocols for report-ing relying party abuse	[?] Isn't relying party certificate + checking against trusted anchor list enough?	No		No	